5. CPU Scheduling

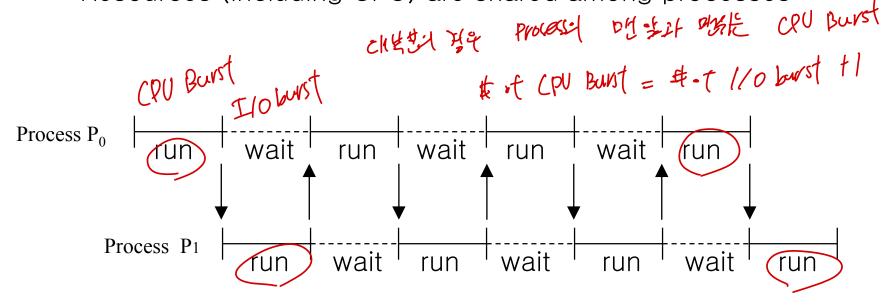
ECE30021/ITP30002 Operating Systems

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Real-time CPU scheduling
- Operation system examples
- Algorithm evaluation

Basic Concepts

- Motivation: maximum CPU utilization obtained with multiprogramming and multitasking
 - Resources (including CPU) are shared among processes



Multiprogramming

CPU-I/O Burst Cycle

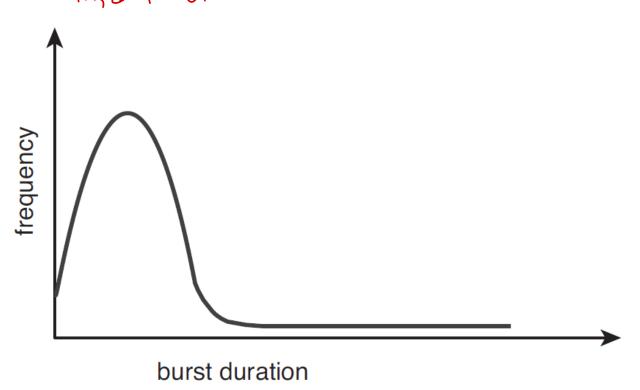
- Process execution consists of cycle of CPU execution and I/O wait
 - First and last bursts are CPU bursts
- Types of processes CPU burst & 1/6 burster
 - I/O-bound process 水和 吃.
 - Consists of many short CPU bursts
 - CPU-bound process
 - Consists of <u>a few long CPU bursts</u>

load store CPU burst add store read from file 1/03 ASK 75 I/O burst wait for I/O store increment index CPU burst write to file I/O burst wait for I/O load store CPU burst add store read from file I/O burst wait for I/O

Histogram of CPU-burst Durations

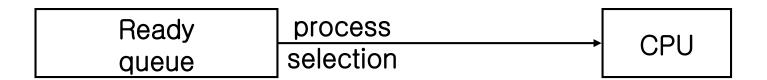
Exponential or hyper exponential

CHEBEL CPV-burter Yolz 36th =) I/o bound process it listed



CPU Scheduler

- CPU scheduler (= short term scheduler)
 - Selects a process from ready queue and allocate CPU to it
- Implementation of ready queue
 - FIFO, priority queue, tree, unordered linked list
 - Each process is represented by PCB



सिन्धिया प्रभ द्वीरं प्रध

Preemptive Scheduling

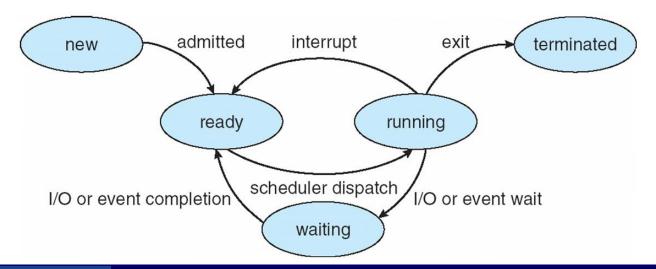
non-preemptive Hts (scheduling)

Preemptive scheduling: scheduling may occur while a process is running.

Ex) interrupt, process with higher priority

When Scheduling Occurs?

- 1. When a process switches from running state to waiting state **
- 2. When a process switches from running state to ready state of town ex) an interrupt occurs
- 3. When a process switches from waiting state to ready state optional ex) completion of I/O 1/0 2 or process from waiting state of process of runing states process of the process of the
- 4. When a process terminates
- -> 1 and 4 are inevitable, but 2 and 3 are optional



Preemptive Scheduling

- Non-preemptive (or cooperative) scheduling
 - Scheduling can occur at 1, 4 only
 - Running process is not interrupted.
- Preemptive scheduling
 - Scheduling can occurs at 1, 2, 3, 4
 - Scheduling may occur while a process is running
 - Requires H/W support and shared data handling

Preemptive scheduling can result in <u>race conditions</u> when data are shared among several processes

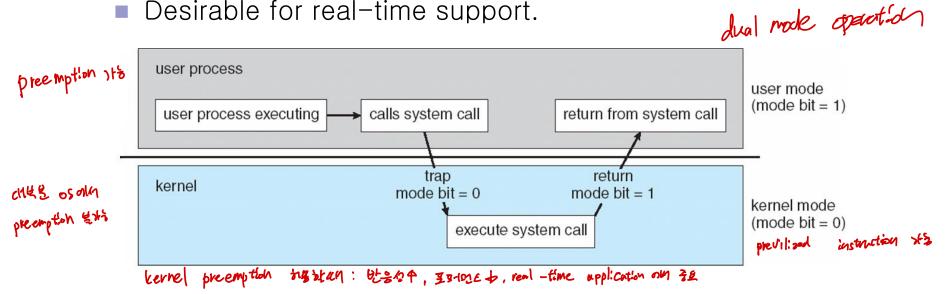
Syncronization of 45

Preemption of OS Kernel

non-preemptur kanel: window, 201 2625

Preemptive kernel kernel allows preemption in system 나 최근 건축소: synchron!zation으로 안장이와 선제 개성

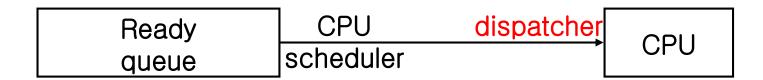
Desirable for real-time support.



- cf. System calls are non-preemptive in most OS.
 - window Keeps kernel structure simple

Dispatcher

- Dispatcher: a module that gives control of CPU to the process selected by short-term scheduler
 - Switching context
 - Switching to user mode
 - Jumping to proper location in user program
- Dispatch latency: time from stopping one process to starting another process



Scheduling Criteria

- CPU utilization: keep the CPU as busy as possible CPUT \$100 to the CPU as busy as possible CPUT \$100 to the CPUT \$100 to t
- Turnaroundtime: interval from submission of a process to its completion
- Waiting time: sum of periods spent waiting in ready queue schedular > tone thing in ready queue
- Response time: time from submission of a request to first response ৬৬%
- -> Importance of each criterion vary with systems
- Measure to optimize
 - Average / minimum / maximum value
 - Variance

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Real-time CPU scheduling
- Operation system examples
- Algorithm evaluation

Scheduling Algorithms

- First-come, first-served (FCFS) scheduling
- Shortest-job-first (SJF) scheduling
- Priority scheduling
- Round-robin scheduling
- Multilevel queue scheduling
- Multiple feedback-queue scheduling

First-Come, First-Served Scheduling



- Process that requests CPU first, is allocated CPU first
 - Non-preemptive scheduling
 - Simplest scheduling method
- Sometimes average waiting time is quite long
 - CPU, I/O utilities are inefficient

Ex) Three processes arrived at time 0

Process	Burst Time	Waiting time	
P1	24	0	
P2	3	24	
P3	2	27	

*time unit: msec

Average waiting time: $(0 + 24 + 27)/3 = 17_{M}$

	P1		P2	P	3
0		24	. 2	27	29

Shortest-Job-First Scheduling

 Assign to the process with the smallest next CPU burst

	Process		Burst Time	Waiting time	
	P1		24	5	Average waiting time:
	P2		3	2	(5 + 2 + 0)/3 = 2.3 ms
	P3		2	0	waiting time optimal
	P3 P2	2		P1	burst time로 영dotifical 정확in 약4가 얼니다
() 2	5			29 b) his by 1893 on 18276

- SJF algorithm is optimal in minimum waiting time
- Problem: difficult to know length of next CPU burst

Shortest-Job-First Scheduling

- Predicting next CPU burst from history
 - Exponential averaging

$$au_{n+1} = lpha t_n + (1-lpha) \,\,\, au_n$$
 ; weighted average

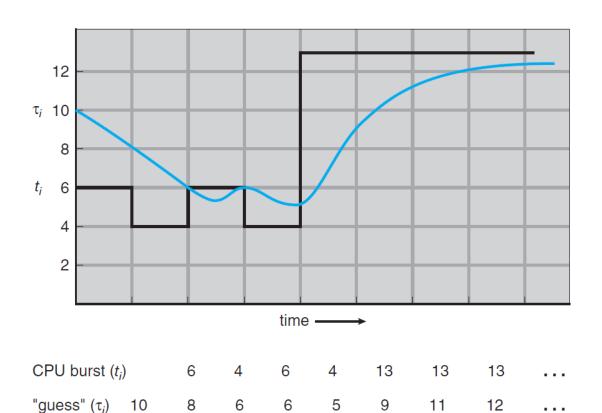
- □ t_n: actual length of n-th CPU burst
- $\square \alpha$: a coefficient between 0 and 1
 - α = 0: recent history has no effect
 - α = 1: only recent history matters
 - \square Usually, $\alpha = 0.5$

Note: older history affects less

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$$

Prediction of CPU Burst

$$\tau_{n+1} = \alpha t_n + (1-\alpha) \tau_n$$



Shortest-Job-First Scheduling



- Shortest-remaining-time-first scheduling (같아된다 - 짜지나) - 식계시간

Process	Arrival Time	Burst Time	Waiting Time
P1	0	780	9
P2	2	AZO	1
P3	4	10	0
P4	5	40	2

0	2	4 5	7	11	 16
(2)	(2)	(1) ((5)	
P1	P2	P3 P	2 P4	P1	

Average waiting time: 3

Priority Scheduling



Equal-priority processes: FCFS

Note: each process has its priority

In this text, lower number means higher priority

外外线线 多鱼

Ex)

Process	Burst Time	Priority	Waiting Time
P1	10	3	6
P2	1	1	0
P3	2	4	16
P4	1	5	18
P5	5	2	1

Average waiting time: 8.2

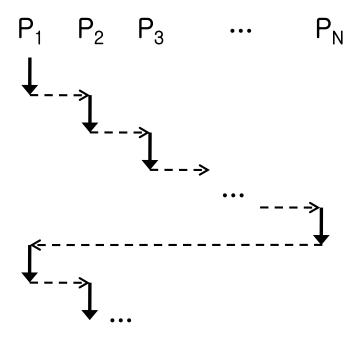
	P2	P5	P1	P3	P4
0	1	(3	16	 18

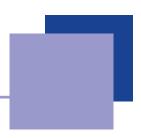
Priority Scheduling

1/0 bound process of could process set

- Priority can be assigned internally and externally
 - Internally: determined by measurable quantity or qualities
 - □ Time limit, memory requirement, # of open files, ratio of I/O burst and CPU burst, ···
 - Externally: importance, political factors
- Priority scheduling can be preemptive or nonpreemptive.
- Major problems
 - Indefinite blocking (= starvation) of processes with lower priorities もならり 火を 芝麻い 水 CPV 七神のた を使とる法
 - -> Solution: aging (gradually increase priority of processes waiting for long time)

- Similar to FCFS, but it's preemptive
 - Designed for time-sharing systems
 - CPU time is divided into <u>time quantum</u> (or <u>time slice</u>)
 - □ A time quantum is 10~100 msec.
 - Cf. switching latency: 10 µsec msc = 1000 msc
 - Ready queue is treaded as circular queue
 - CPU scheduler goes around the ready queue and allocate
 CPU time up to 1 time quantum

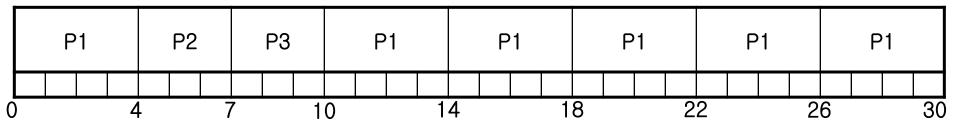




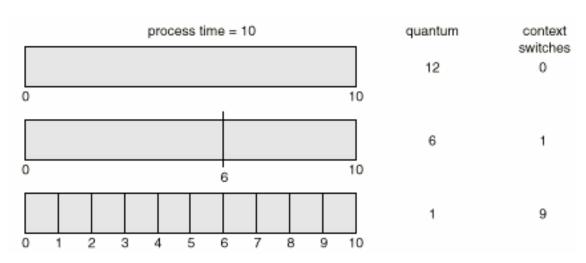
Ex) Time quantum = 4

Process	Burst Time	Waiting Time	
P1	24	6	
P2	3	4	
P3	3	7	

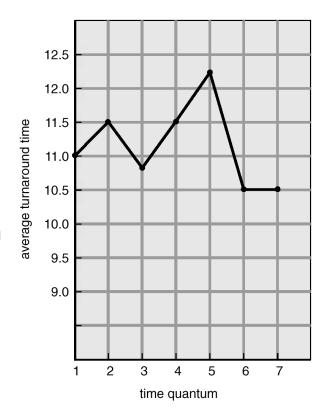
Average waiting time: 5.66



- Performance of RR scheduling heavily depends on size of time quantum.
 - Time quantum is small: processor sharing
 - Time quantum is large: FCFS
- Context switching overhead depends on size of time quantum.



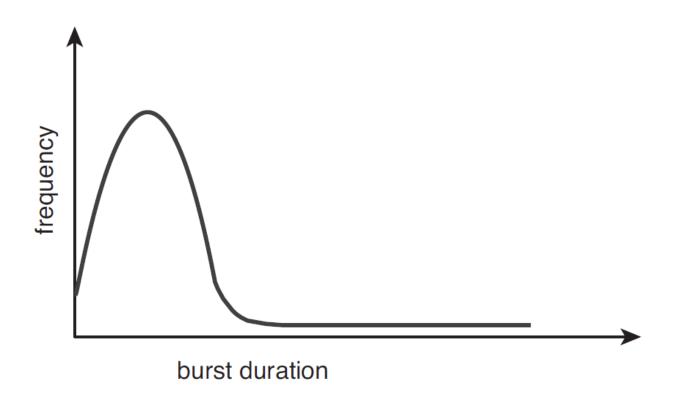
- Turnaroundtime also depends on size of time quantum
 - Average turnaround time is not proportional nor inverse-proportional to size of time quantum
 - Average turnaroundtime is improved if most processes finish their next CPU burst in a single time quantum
 - However, too long time quantum is not desirable
 - A rule of thumb: about 80% of CPU burst should be shorter than time quantum



process	time
P ₁ P ₂ P ₃ P ₄	6 3 1 7

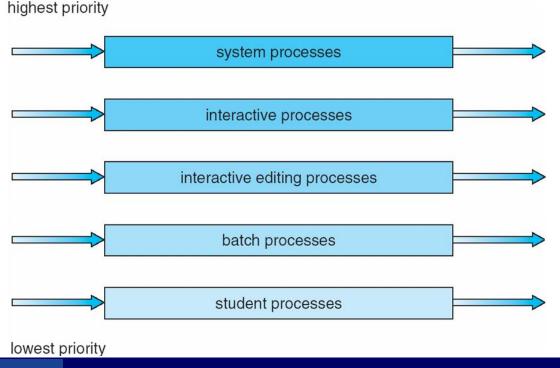
Histogram of CPU-burst Durations

Exponential or hyper exponential



Multilevel Queue Scheduling

- Classify processes into different groups and apply different scheduling
 - Memory requirement, priority, process type, …
- Partition ready queue into several separate queues



Multilevel Queue Scheduling

- Each queue has its own scheduling algorithm
- Scheduling among queues
 - Fixed-priority preemptive scheduling
 - A process in lower priority queue can run only when all of higher priority queues all empty
 - Time-slicing among queues
 - Ex) foreground queue (interactive processes): 80% background queue (batch processes): 20%
- Assignment of a queue to a process is permanent.

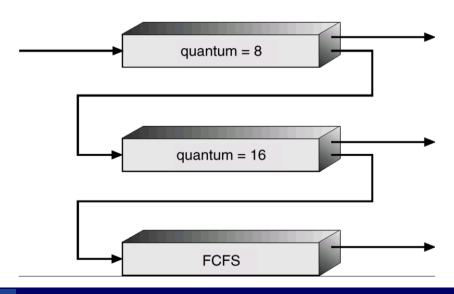


Multilevel Feedback-Queue Scheduling

- Similar to multilevel queue scheduling, but a process can move between queues.
- Idea: separate processes according to characteristics of their CPU bursts.
 - If a process uses too much CPU time, move it to lower priority queue.
 - I/O-bound, interactive processes are in higher priority queues.

Multilevel Feedback-Queue Scheduling

- Ex) Ready queue consists of three queues (0~2)
 - Q_0 (time limit = 8 milliseconds)
 - Q_1 (time limit = 16 milliseconds)
 - Q_2 (FCFS)
- → A new process is put in Q₀. if it exceeds time limit, it moves to lower priority queue



Multilevel Feedback-Queue Scheduling

- Parameters to define a multilevel feedback-queue scheduler
 - # of queues
 - Scheduling algorithm for each queue
 - Method to determine when to upgrade a process to higher priority queue
 - Method to determine when to demote a process to lower priority queue
 - Method to determine which queue a process will enter when it needs service
- → The most complex algorithm

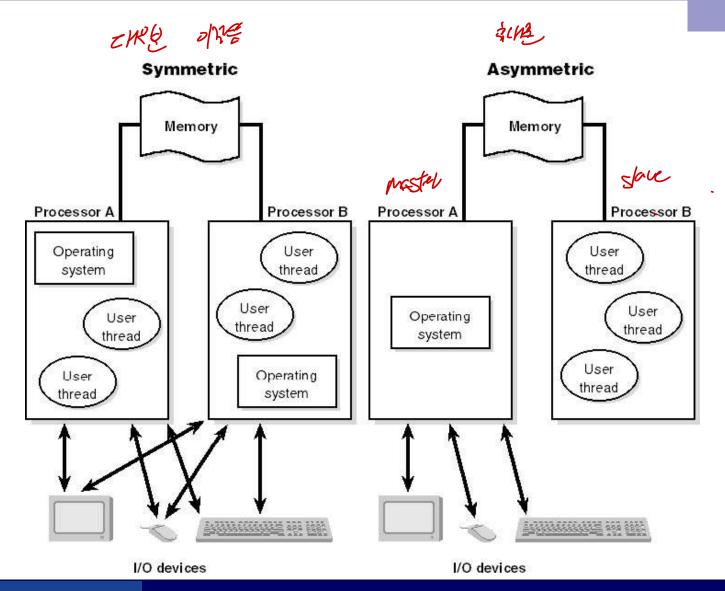
Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Real-time CPU scheduling
- Operation system examples
- Algorithm evaluation

Multiple-Processor Scheduling

- Multiple-processor system
 - Load sharing is possible.
 - Scheduling problem is more complex.
- There are many trials in multiple-processor scheduling, but no generally best solution
- In this text, all processors are assumed identical.
 - Any process can run on any processor

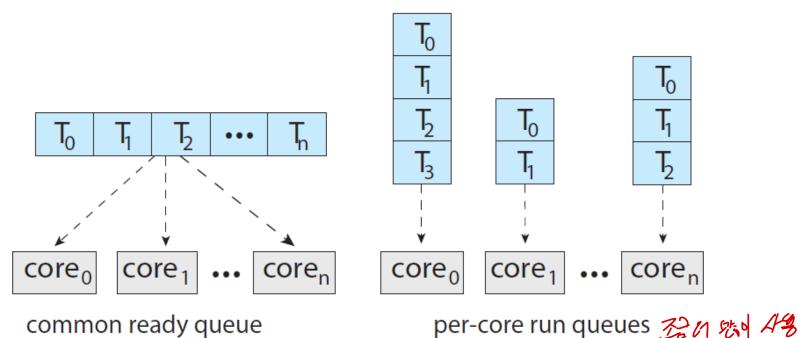
Symmetric vs. Asymmetric Multiprocessing



Multi-processor Scheduling

Two possible strategies

- 1. All threads may be in a common ready queue.
 - Needs to ensure two processors do not choose the same thread nor lost from the queue
- 2. Each processor may have its own private queue of threads.

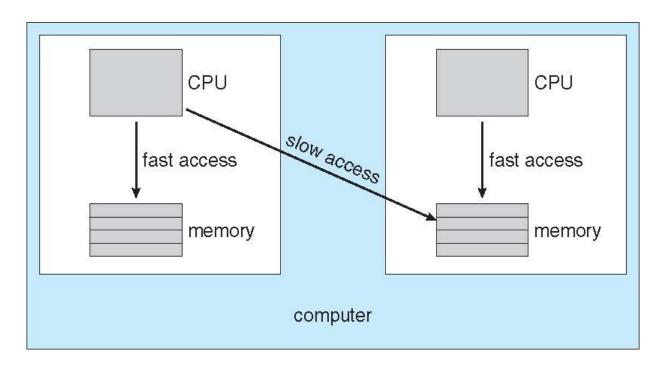


Processor Affinity

- Overhead of migration of processes from one processor to another processor
 - All contents of cache should be invalidated and repopulated
- Processor affinity: keeping a process running on the same processor to avoid migration overhead
 - Soft affinity: Although OS attempts to keep a process running on the same processor, a process may migrate between processors.
 - Hard affinity: It is possible to prevent a process from being migrated to other processor

Processor Affinity

- NUMA(Non-Uniform Memory Access) and CPU scheduling
 - A CPU has faster access to some parts of main memory than to other parts



Load Balancing

- Load balancing: attempt to keep the workload evenly distributed across all processors.
 - Necessary for system where each processor has its own ready queue (It's true for most modern OS's)
- Two general approaches
 - Push migration
 - A specific task periodically checks the load on each processor
 - ☐ If unbalance is found, move processes to idle or less-busy processors
 - Pull migration
 - ☐ An idle processor pulls a waiting task

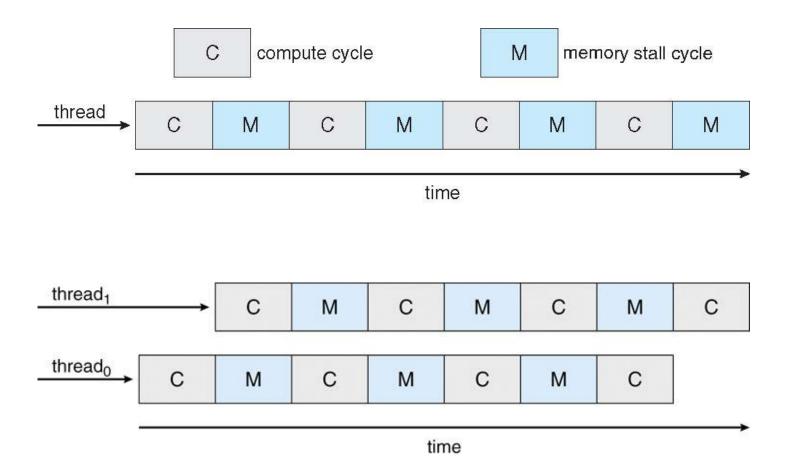
Note! Push and pull migration can be implemented in parallel.

- -> Linux, ULE scheduler for FreeBSD
- Load balancing can counteract the benefit of processor affinity

Multi-core Processors

- Multi-core processor: multiple processor cores on the same physical chip
 - Recent trend
 - Faster and consume less power than systems in which each processors has its own physical chip
- Scheduling issues on multi-core processor
 - Memory stall: for various reasons, memory access spends significant amount time waiting for the data to become available.
 - Ex) accessing data not in cache
 - Remedy: multithreaded processor cores
 - □ Two or more hardware threads are assigned to each core

Multi-Threaded Processor Cores

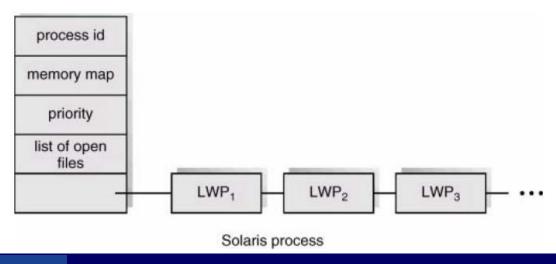


Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Real-time CPU scheduling
- Operation system examples
- Algorithm evaluation

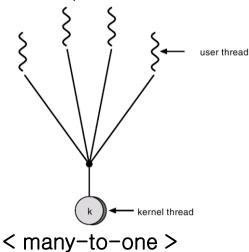
Thread Scheduling

- Threads
 - User thread -> supported by thread library
 - Scheduled indirectly through LWP
 - Kernel thread -> supported by OS kernel
- Actually, it is not processes but kernel threads that are being scheduled by OS



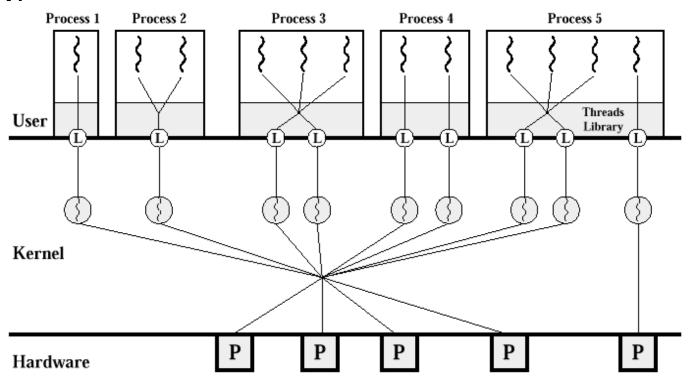
Contention Scope

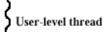
- Process-contention scope (PCS)
 - Competition for LWP among user threads
 - Many-to-one model or many-to-many model
 - Priority based
- System-contention scope (SCS)
 - Competition for CPU among kernel threads



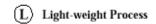
Scheduler Activation and LWP

Connection between user / kernel threads through LWP











Pthread Scheduling

- In thread creation with Pthreads, we can specify PCS or SCS
 - PCS(PTHREAD_SCOPE_PROCESS):
 - □ In many-to-one, only PCS is possible
 - SCS(PTHREAD_SCOPE_SYSTEM):
 - □ In one-to-one model, only SCS is possible

Note: On certain system, only certain values are allowed

- □ Linux, MacOS X
- Related functions
 - pthread_attr_setscope(pthread_attr_t *attr, int scope)
 - pthread_attr_getscope(pthread_attr_t *attr, int *scope)

Pthread Scheduling

Example)

```
pthread t tid;
pthread_attr_t attr;
int scope;
pthread_attr_init(&attr);
pthread_attr_getscope(&attr, &scope);
// scope is either PTHREAD_SCOPE_PROCESS
// or PTHREAD_SCOPE_SYSTEM
pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM)
pthread_create(&tid, &attr, runner, NULL);
```

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Real-time CPU scheduling
- Operation system examples
- Algorithm evaluation

Real-time CPU Scheduling 组 水 如此 咖啡 。

Revi-time Option: 장하인 Etital이 시캠드 보상하는 시스템 appliacion + as

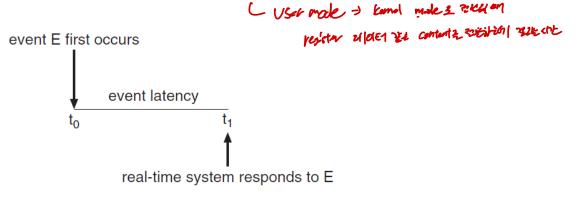
- Real-time Operating Systems (RTOS)
 - OS intended to serve real-time systems
 Ex) Antirock brake system (ABS) requires latency of 3~5 msec
- Soft real-time systems
 - No guarantee as to when a critical real-time process will be scheduled.
 - Guarantee only that the process will be given preference over noncritical processes.
- - A task must be serviced by its deadline
 - Service after the deadline has expired is the same as no service at all.

Minimizing Latency

Hill Well It Systems It evand-driven system

- Most real-time systems are event-driven system
 - When an event occurs, the system must respond to and service it as quickly as possible
- Event latency: the amount of time that elapses from when an event occurs to when it is serviced
 - Interrupt latency, dispatch latency

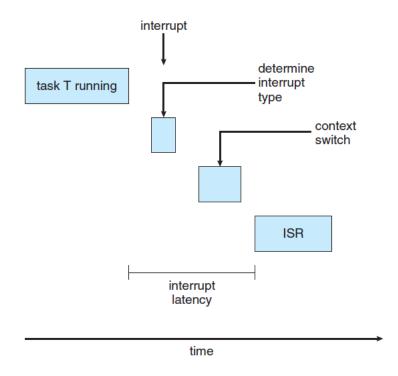
interrupt it were interrupt interrup



Time

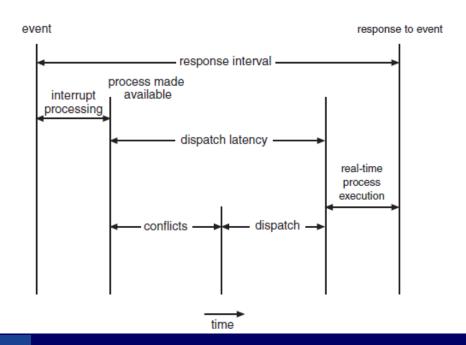
Interrupt Latency

Interrupt latency: the period of time from the arrival of interrupt at the CPU to the start of the routine that services the interrupt.



Dispatch Latency

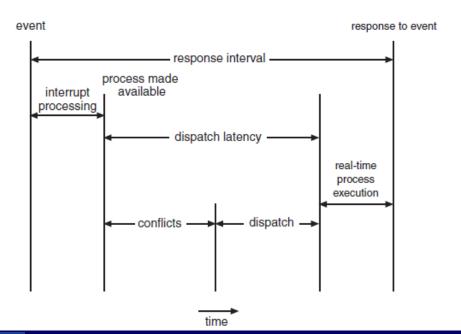
- Dispatch latency: the amount of time required for the scheduling dispatcher to stop one process and start another
 - Preemptive kernels are the most effective technique to keep dispatch latency low.



Dispatch Latency

By Enorth processif about spice, the processif states exten

- Conflict phase
 - 1. Preemption of any process running in the kernel
 - 2. Release of resources occupied by low-priority processes needed by a high-priority process
- Dispatch phase schedules the high-priority process onto an available CPU.



Priority-based Scheduling

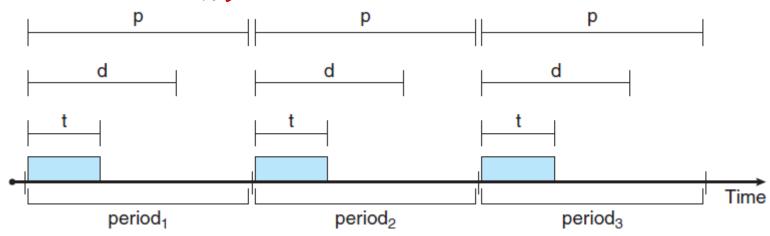
- The most important feature of RTOS is to respond immediately to a real-time process
 - The scheduler should support priority-based preemptive algorithm
 - Most OS assign highest priority to real-time processes
 - Preemptive, priority-based scheduler only guarantees soft real-time functionality.
- Hard real-time systems must further guarantee that real-time tasks will be serviced in accord with their deadline requirements
 - Requires additional scheduling features.

Periodic Process 218 and 30 mil 32 mi

- - A fixed processing time t
 - A deadline d by which it must be serviced by the CPU
 - A period p

$$0 \le t \le d \le p$$

■ The rate of a periodic task: 1/p



Scheduling with Deadline Requirement

 A process may have to announce its deadline requirements to the scheduler

```
Shit size processit dad [in on int itsize ite
```

- Using an admission-control algorithm, the scheduler does one of two things:
 - Admits the process, guaranteeing that the process will complete on time, or
 - Rejects the request as impossible if it cannot guarantee that the task will be serviced by its deadline.

Rate-Monotonic Scheduling

- ZWE THE PROPERTY ZWE
- Rate-monotonic scheduling algorithm
 - Schedules periodic tasks using a static priority policy with preemption.
 - Each periodic task is assigned a priority inversely based on its period. periodic task is assigned a priority inversely based on its period. periodic task is assigned a priority inversely based on its period. periodic task is assigned a priority inversely based on
 - The processing time of a periodic process is assumed the same for each CPU burst.

Rate-Monotonic Scheduling



Rate-monotonic scheduling is considered optimal

If a set of processes cannot be scheduled by this algorithm, it cannot be scheduled by any other algorithm that assigns static priorities.

Limitation

- <u>CPU utilization</u> is bounded, and it is not always possible to maximize CPU resources fully.
- Worst-case CPU utilization: $N(2^{1/N} 1)$
 - \square N: # of processes

Earliest-Deadline-First Scheduling

periodic 741 oboth 55

- Earliest-deadline-first (EDF) scheduling
 - The earlier the deadline, the higher the priority

Requirements

- When a process becomes runnable, it must announce its deadline requirements to the system.
- Does not require that processes be periodic, nor must a process require a constant amount of CPU time per burst.

Theoretically optimal

Theoretically, it can schedule processes so that each process can meet its <u>deadline requirements</u> and <u>CPU</u> <u>utilization</u> will be 100 percent.

Proportional Share Scheduling



- Allocates T shares among all applications.
- An application can receive N shares of time, thus ensuring that the application will have N / T of the total processor time
- Must work in conjunction with an admission-control policy to guarantee that an application receives its allocated shares of time



POSIX Real-Time Scheduling

- Scheduling classes for real-time threads
 - SCHED_FIFO: FCFS policy
 - □ No time slicing among threads of equal priority
 - The highest-priority real-time thread at the front of the FIFO queue will be granted the CPU until it terminates or blocks
 - SCHED_RR: a round-robin policy.
 - □ Similar to SCHED_FIFO
 - □ Time slicing among threads of equal priority.
 - SCHED_OTHER: system-specific
 - Implementation is undefined
- POSIX API for getting/setting scheduling policy
 - pthread_attr getschedpolicy(pthread_attr_t *attr, int *policy);
 - pthread_attr setschedpolicy(pthread_attr_t *attr, int policy);

POSIX Real-Time Scheduling

```
int i = 0, policy = 0;
pthread_t tid[NUM_THREADS];
pthread_attr_t attr;
/* get the default attributes */
pthread_attr_init(&attr);
/* get the current scheduling policy */
if(pthread_attr_getschedpolicy(&attr, &policy)
!= 0)
   fprintf(stderr, "Unable to get policy.₩n");
else {
   if(policy == SCHED_OTHER)
      printf("SCHED OTHER₩n");
   else if(policy == SCHED_RR)
      printf("SCHED RR₩n");
   else if(policy == SCHED_FIFO)
      printf("SCHED FIFO\(\fomath{\psi}\)n");
```

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Real-time CPU scheduling
- Operation system examples
- Algorithm evaluation

1 37,47 W block

4) 37~45 (ine ladle) 34:

: 37/45 V/ block 5 2571 Writh 7

: p block
[(alc n/2/22 489)]

Operating System Examples

- Linux
- Windows
- (Solaris)

Linux Scheduler

History

- Traditional UNIX scheduling algorithm (ver. < 2.5)
 - Not adequate support for SMP system
 - Not scale well as # of tasks increases
- O(1) scheduling algorithm (ver. ≥ 2.5)
 - □ O(1) complexity
 - □ Increased support for SMP
 - □ Poor response time interactive processes
- Completely Fair Scheduler (CFS) (ver. ≥ 2.6)

Linux Scheduler

Based on scheduling classes

- Each class is assigned a specific priority.
 - Default scheduling class (CFS algorithm)
 - □ Static priority 100 ~ 139
 - □ Real-time scheduling class (SCHED_FIFO, SCHED_RR)
 - □ Static priority 0 ~ 99
 - □ If necessary, new scheduling classes can be added
- Allows different scheduling algorithms based on the needs of the system (e.g. server vs. mobile)
- Scheduler selects highest-priority task belonging to highest-priority class

Linux CFS Scheduler

- CFS scheduler assigns a proportion of CPU processing time to each task
 - The portion is calculated from nice value.
 - □ Nice value: relative priority in [-20, +19], default is 0
 - ☐ Higher value represents lower priority ('nice' to other tasks)
 - Proportions of CPU time are allocated from the value of targeted latency.
 - Targeted latency: interval of time during which every runnable task should run at least once
 - Default and minimum values
 - Can increase if the number of active tasks in the system grows beyond a certain threshold

Linux CFS Scheduler

- The scheduler simply selects the task that has the smallest vruntime value
 - Per-task variable vruntime(virtual run time)
 - vruntime = <actual physical run time> * <decay factor>
 - Decay factor
 - □ Normal priority tasks: 1.0
 - □ Lower priority task: higher decay factor (> 1.0)
 - □ Higher priority task: lower decay factor (< 1.0)</p>
 - → Higher priority tasks are likely to have smaller vruntime
- Higher-priority task can preempt a lower-priority task.

Linux CFS Scheduler

- I/O-bound tasks: run only for short periods before blocking for additional I/O
 - → vruntime will be lower
- CPU-bound tasks: exhaust its time period whenever it has an opportunity to run on a processor
 - vruntime will be higher
- → Gives I/O-bound tasks higher priority than CPU-bound tasks

Windows Scheduling



- Priority-based, preemptive scheduling
 - Dispatcher handles scheduling
 - Dispatcher ensures the highest priority thread will always run
 - 32-level priority scheme
 - □ Real-time class (16~32)
 - □ Variable class (1~15)
 - Memory management (0)
 - If no ready thread is found, dispatcher runs idle threads

Windows Scheduling

Priority classes

- A process belongs to one of 6 levels of priority class
- Each thread belongs to one of 7 levels of relative priority

	real- time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

Windows Scheduling

Altering priority

- Priority class of a process: SetPriorityClass()
- Base priority of a thread: SetThreadPriority()

Adjusting priority for variable-priority class

- Lower priority of a thread that expired entire time quantum
 - □ Never lowered below the base priority
- Boost priority of a thread that released from waiting
 - Amount of increment vary with what it was waiting

Scheduling of foreground / background processes

 Time quantum for the foreground process is increased by some factor (typically, 3)

Agenda

- Basic concepts
- Scheduling algorithms
- Multiple-processor scheduling
- Thread scheduling
- Real-time CPU scheduling
- Operation system examples
- Algorithm evaluation

- First of all, evaluation criteria should be determined Ex)
 - 1. Maximum response time is 1 second
 - 2. Satisfying 1, maximize CPU utilization
 - 3. Maximize throughput such that turnaround time is linearly proportional to total execution time

Evaluation methods

- Deterministic modeling
- Queueing models
- Simulations

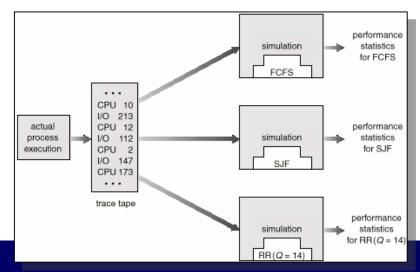
Deterministic modeling

Evaluation for a particular predetermined workload

Queueing models

- Evaluation based on probabilistic distribution of...
 - □ CPU and I/O burst
 - Arrival time of CPU burst
- Description of system
 - □ Network of servers, each of which has a ready queue
 - Given arrival rates and service rates, performance can be computed probabilistically

- Simulations, using a programming model of computer system
 - Generator of process, CPU burst time, arrival, departure, …
 - Usually, the generator randomly generates the simulated events based of some distribution
 - Limitation: simulated situation isn't exactly same with the real situation
 - □ Especially, distribution gives no information about order of events
 - Remedy: using a record of real system. (trace tape)



- Implementation: the only way for accurate evaluation
 - Performance depend not only on scheduling algorithm and OS support, but also user's interaction
 - Environment will change
 - Ex) If short process is given higher priority, user may break a large process into several smaller processes.
 - Some OS's provides flexibility to alter scheduling scheme
 - Performance turning for specific (set of) applications
 - API's to modify priority of a process or thread