

# 6. Process *(thread)* Synchronization

ECE30021/ITP30002 Operating Systems

# Agenda

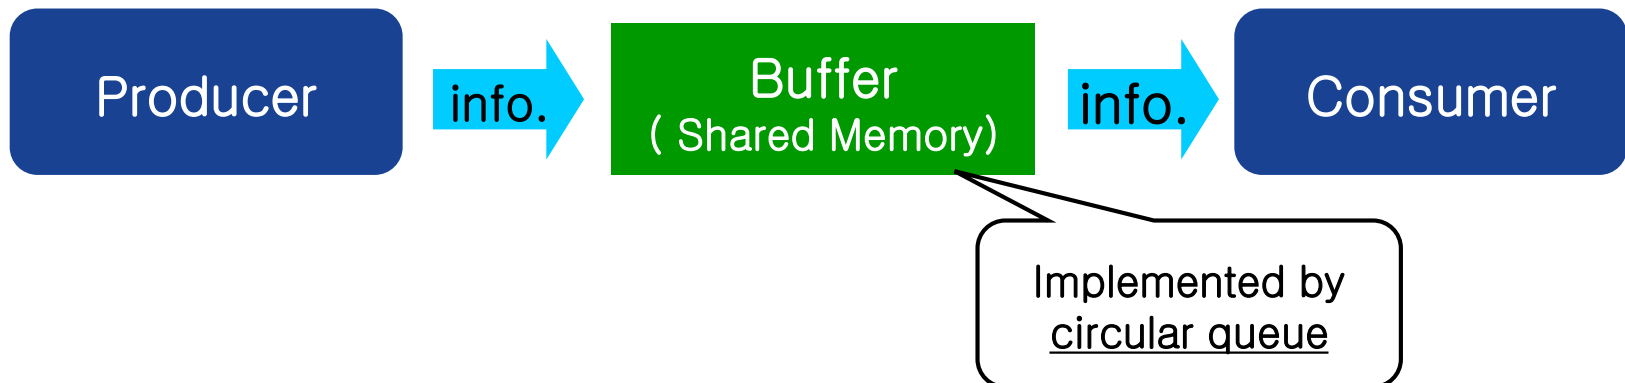
---



- Background
- The critical-section problem
- H/W Support for Synchronization
- Mutex and Semaphore
- Monitors
- Liveness

# Background

- Process communication method
  - Message passing
  - Shared memory → confliction can occur !!
- Producer–consumer problem
  - An example of communication through shared memory



# Concurrent Access of Shared Data

## ■ Producer *que insertion*

```
while(true){  
    while(counter==BUFFER_SIZE);  
    buffer[in] = nextProduced;  
    in = (in+1)%BUFFERSIZE;  
    counter++;  
}
```

*한개 item의 갯수*

## ■ Consumer

```
while(true){  
    while(counter==0);  
    nextConsumed = buffer[out];  
    out = (out+1)%BUFFER_SIZE;  
    counter--;  
}
```

## ■ Implementation of “counter++”

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

## ■ Implementation of “counter--”

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

Switching can occur during modification operations !

# Synchronization Problem

## ■ Problematic situation

- Initial value of counter is 5.
- Producer increased counter, and consumer decreased counter concurrently.
- Ideally, counter should be 5. But...

Producer	Consumer
$\text{register}_1 = \text{counter}$ ( <i>register1 = 5</i> ) $\text{register}_1 = \text{register}_1 + 1$ ( <i>register1 = 6</i> )  $\text{counter} = \text{register}_1$ ( <i>counter = 6</i> )	$\text{register}_2 = \text{counter}$ ( <i>register2 = 5</i> ) $\text{register}_2 = \text{register}_2 - 1$ ( <i>register2 = 4</i> )  $\text{counter} = \text{register}_2$ ( <i>counter = 4</i> )

counter can be 4 or 6 !

# Process Synchronization



## ■ Race condition

- A situation, where several processes access and manipulate the same data concurrently
- The outcome of execution depends on the particular order in which the access takes place.

## ■ Synchronization: the coordination of occurrences to operate in unison with respect to time.

Ex) ensuring only one process can access the shared data at a time

# Agenda

---



- Background
- The critical-section problem
- H/W Support for Synchronization
- Mutex and Semaphore
- Monitors
- Liveness

# The Critical-Section Problem

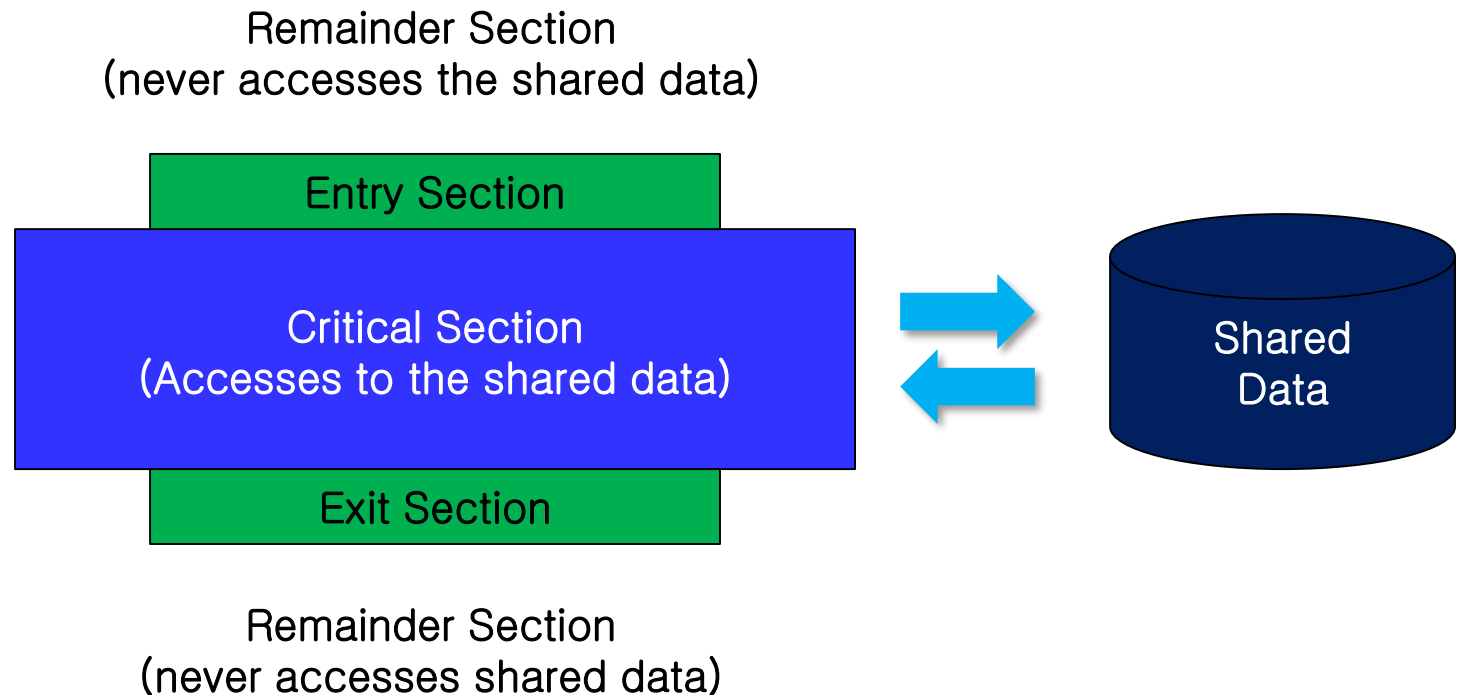
- **Critical section problem**: designing a protocol that processes can use to cooperate
  - $n$  processes  $P_0, \dots, P_{n-1}$  are running on a system concurrently. Each process want access a shared data
  - The code of each process is composed of
    - **Critical section**: a segment of code which may change shared resources (common var., file, ...)
    - **Remainder section**: a segment of code which doesn't change shared resources
    - **Entry section**: code section to request permission to enter critical section
    - **Exit section**: code section to notice it leaves the critical section

race condition  
eg: 1

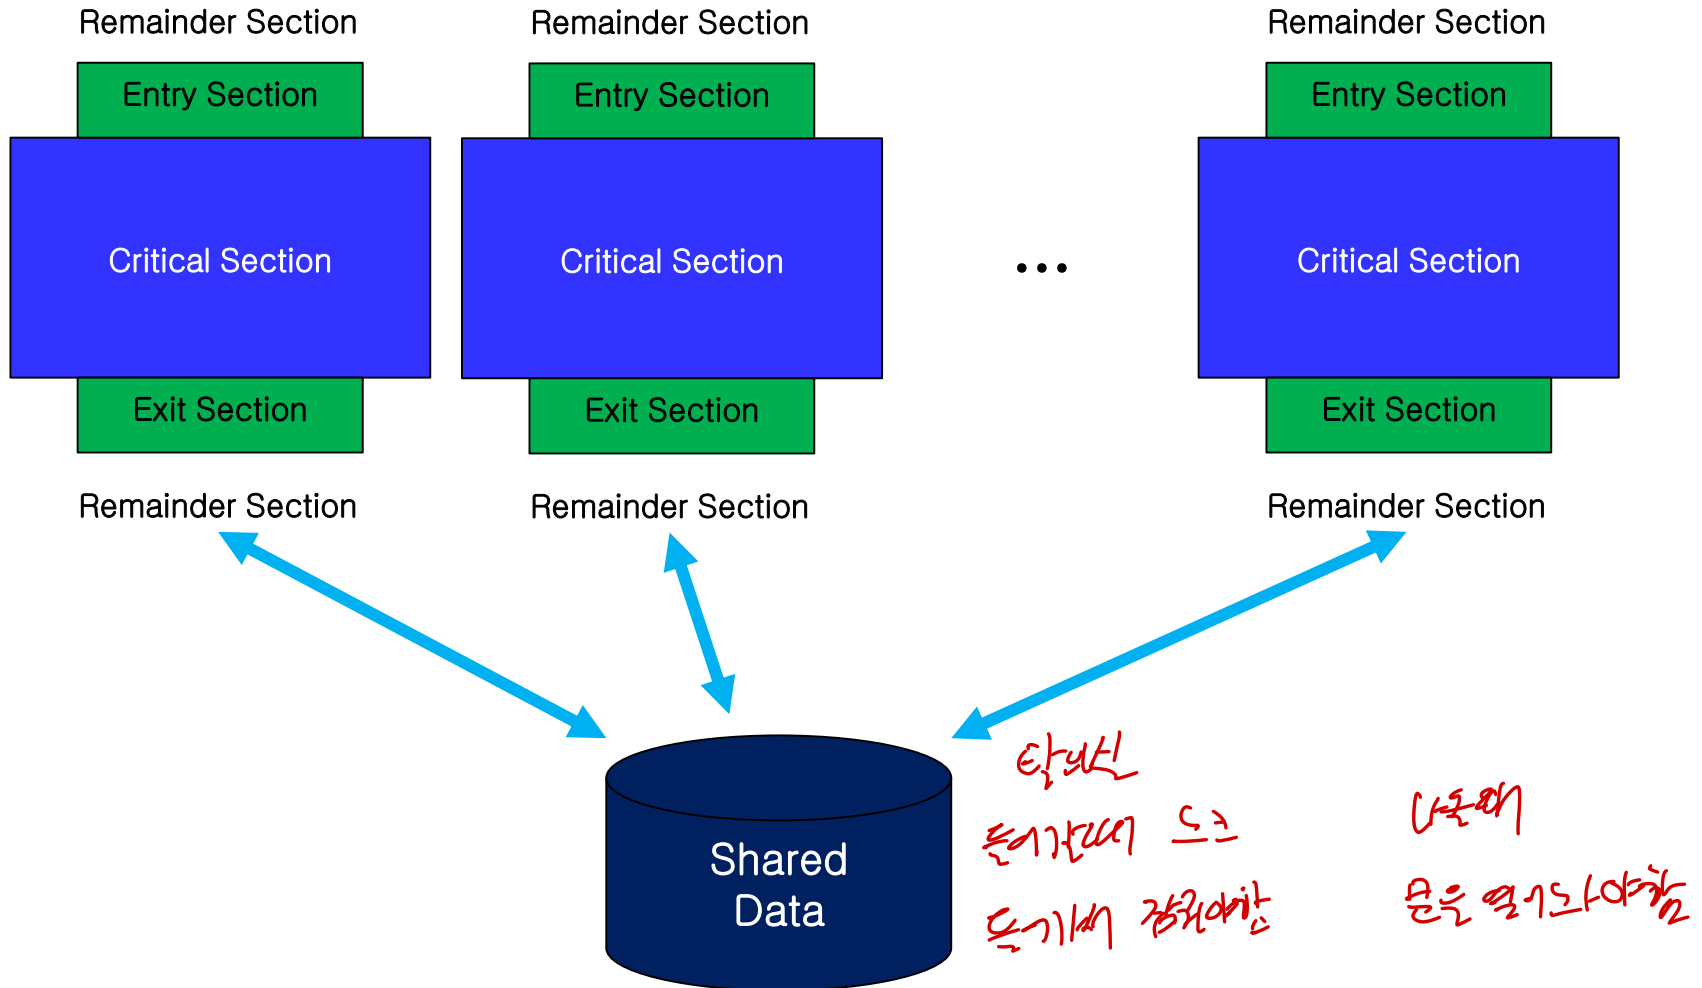


# Critical Section

- A process <sup>may access and</sup> modifies shared data only in the critical section
- At a time only one process can exist in its critical section.



# Critical Section



# The Critical-Section Problem

- General structure of typical processes

do {

*entry section*

critical section

*exit section*

remainder section

} while(TRUE);

# The Critical-Section Problem

## ■ Requirements of critical-section problem

- **Mutual exclusion**: If a process is in its critical section, no other processes can be executing in their critical sections  
→ Other processes should wait *entry section에서 기다림*
- **Progress**: If no process is executing in its critical section, and some processes wish to enter their critical sections, only processes not executing in their remainder section can participate in the decision of next process to enter its critical section next. This selection cannot be postponed indefinitely.  
*누가 작업할 것인지 상황이 아니라 어떤 과정을 먼저 할지 결정한다.*
- **Bounded waiting**: There exists a bound or limit on number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before it is granted.

# The Critical-Section Problem



- Kernel is an example of critical-section problem
  - Kernel data structures such as list of open files,
- Approaches
  - **Non-preemptive kernel**: switching cannot occur when a process is executing in kernel mode
    - Free from race condition
    - Ex) Windows 2000, Windows XP, early version of UNIX, Linux prior to v. 2.6.
  - **Preemptive kernel** accompanied with a solution of critical-section problem
    - Suitable for real-time programming
    - More responsive
    - Ex) Linux v. 2.6, Solaris, IRIX

# Peterson's Solution

- A S/W solution for critical section problem

- ➔ No guarantee to work correctly on some architectures due to **load/store** instructions, but helps to understand the problem

- Two processes  $P_0$  and  $P_1$  (or  $P_i$  and  $P_j$ )  
*Current process* *other process*

- Data items

```
int turn;           // indicates process allowed to execute
                    // in its critical section
                    // initial value is 0
```

```
boolean flag[2]; // if flag[i] is true,  $P_i$  is ready to enter
                 // its critical section
```

# Erroneous Algorithm 1

## ■ Algorithm 1

### ■ Process $P_i$

do {

    while (turn  $\neq$  i);

        critical section

    turn = j;

        remainder section

} while (1);

# Erroneous Algorithm 1

progress ~~ex~~

## ■ Process $P_0$

```
do {  
    while (turn != 0) ;  
        critical section  
    turn = 1;  
        remainder section  
} while (1);
```

## ■ Process $P_1$

```
do {  
    while (turn != 1) ;  
        critical section  
    turn = 0;  
        remainder section  
} while (1);
```

## ■ Mutual exclusion is guaranteed, but progress is not guaranteed.

Ex)  $P_1$  is trying to enter its critical section, but  $P_0$  is in its remainder section  $\rightarrow$  turn is not switched to 1



# Erroneous Algorithm 2

## ■ Algorithm 2

### ■ Process $P_i$

do {

    flag[i] = true;

    while (flag[j]) ;

        critical section

    flag[i] = false;

        remainder section

} while (1);

# Erroneous Algorithm 2

progress ~~is~~ X

## ■ Process $P_0$

```
do {  
    flag[0] = true;  
    while(flag[1]);  
    critical section  
    flag[0] = false;  
    remainder section  
} while (1);
```

## ■ Process $P_1$

```
do {  
    flag[1] = true;  
    while(flag[0]);  
    critical section  
    flag[1] = false;  
    remainder section  
} while (1);
```

■ Mutual exclusion is guaranteed, but progress is not guaranteed.

Ex)  $P_0$  and  $P_1$  enter simultaneously. Both of flag[0] and flag[1] can be true

# Peterson's Solution



## ■ Peterson's Solution

■ Process  $P_i$

```
do {  
    flag[i] = true;  
    turn = i;  
    while (flag[j] and turn == j);  
        critical section  
    flag[i] = false;  
        remainder section  
} while (1);
```

■ Process  $P_j$

```
do {  
    flag[j] = true;  
    turn = j;  
    while (flag[i] and turn == i);  
        critical section  
    flag[j] = false;  
        remainder section  
} while (1);
```

## ■ Satisfies the three conditions

# Peterson's Solution



## ■ Peterson's Solution

■ Process  $P_0$

```
do {  
    flag[0] = true;  
    turn = 1;  
    while (flag[1] and turn == 1);  
        critical section  
    flag[0] = false;  
        remainder section  
} while (1);
```

■ Process  $P_1$

```
do {  
    flag[1] = true;  
    turn = 0;  
    while (flag[0] and turn == 0);  
        critical section  
    flag[1] = false;  
        remainder section  
} while (1);
```

## ■ Satisfies the three conditions

# Peterson's Solution

## ■ Proof of mutual exclusion

- If both  $P_i$  and  $P_j$  enter their critical section, it means
  - $\text{flag}[0] = \text{flag}[1] = \text{true}$
  - $\text{turn}$  can be either 0 or 1, but  $\text{turn}$  cannot be both

## ■ Proof of progress and bounded waiting

- Blocking condition of  $P_i$ :  $\text{flag}[j] == \text{true}$  and  $\text{turn} == j$ 
  - If  $P_j$  is not ready to enter critical section:  $\text{flag}[j] == \text{false}$   
→  $P_i$  can enter critical section
  - If  $P_j$  waiting in its while statement,  $\text{turn}$  is either  $i$  or  $j$
  - If  $\text{turn}$  is  $i$ ,  $P_i$  will enter critical section.
  - Otherwise,  $P_j$  will enter critical section.
  - When  $P_j$  exits critical section,  $P_j$  sets  $\text{flag}[j]$  to false and  $P_i$  can enter critical section because  $P_j$  modifies  $\text{turn}$  to  $i$ .
  - Therefore,  $P_i$  will enter critical section (Progress) and waits at most one process (Bounded waiting).

# Peterson's Solution

process가 2개일때만 가능

- Not guaranteed to work on modern computer architectures
  - To improve system performance, processors and/or compilers may reorder read and write operations that have no dependencies. 정리하면 순서가 바뀌어 실행
  - For a multithreaded application with shared data, the reordering of instructions may render inconsistent or unexpected results.

# Peterson's Solution

## ■ Example

### ■ Shared variables

boolean flag = false;

int x = 0;

### ■ Thread1

1. while(!flag);

2. print x;

### ■ Thread2

1. x = 100;

2. flag = true;

→ No dependency between line 1 and 2 → can be reordered

*CPU  
정확성* ) or optimization reason이 순서를 바꿔 둘 수 있음

# Agenda

---



- Background
- The critical-section problem
- H/W Support for Synchronization
- Mutex and Semaphores
- Monitors
- Liveness



# Memory Model



- **Memory model:** how a computer architecture determines what memory guarantees it will provide to an application program.
- Categories of memory models
  - **Strongly ordered:** a memory modification on one processor is immediately visible to all other processors
  - **Weakly ordered:** modifications to memory on one processor may not be immediately visible to other processors

# Memory Barriers

low-level instruction

kernel(커널)에서의 사용

- **Memory barrier** (or **memory fences**): instructions that can *force* any changes in memory to be propagated to all other processors  
전파하다
  - All loads and stores are completed before any subsequent load or store operations are performed
- **Example**
  - Thread1

```
while (!flag)
    memory barrier();
print x;
```
  - Thread2

```
x = 100;
memory barrier();
flag = true;
```

# Hardware Instructions

- Critical section problem requires lock

- Race condition can be prevented by protecting critical section by **lock**

do {

*acquire lock*

critical section

*release lock*

remainder section

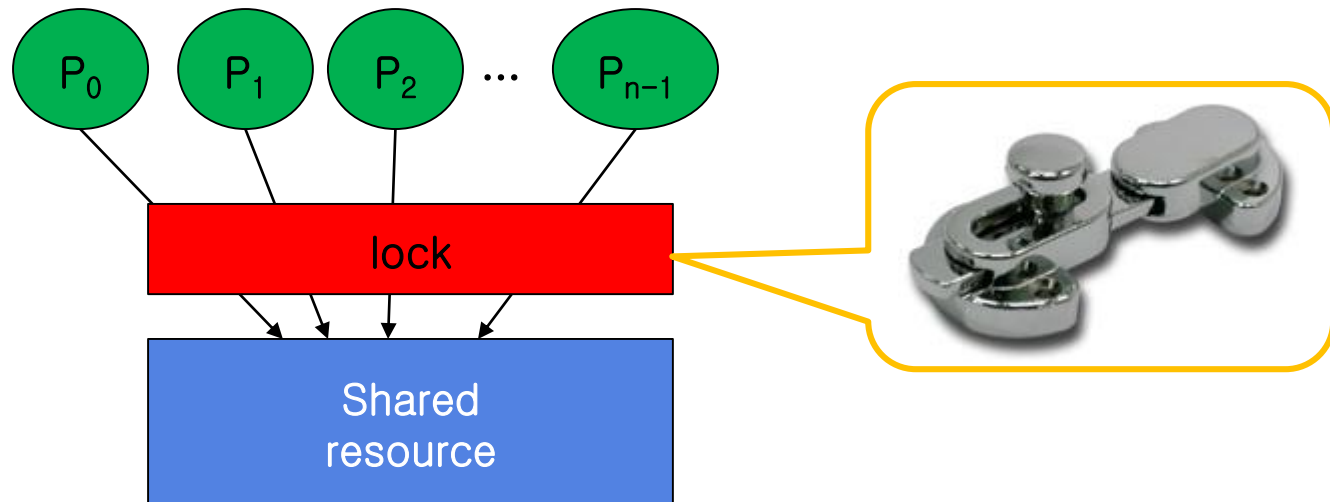
} while(TRUE);

- H/W support makes it easier and improve efficiency

# Mutual Exclusion by Lock variable

## ■ A shared variable *lock*

- Boolean *lock* = 0;
- If *lock* is false, any process can enter its critical section
  - When a process enters its critical section, it sets *lock* to true
- If *lock* is true, a process ( $P_i$ ) is in its critical section.
  - Other processes should wait until  $P_i$  leaves its critical section and sets *lock* to false



# Mutual Exclusion by Lock variable

- Initially, lock = false

- $P_0$   
do {  
    while(lock);  
    lock = true;  
    critical section  
    lock = false;  
    remainder section  
} while(TRUE);

*knocking*

*locking*

- $P_1$   
do {  
    while(lock);  
    lock = true;  
    critical section  
    lock = false;  
    remainder section  
} while(TRUE);

- Note! Checking and locking must not be separated

# Interrupt Disable



- In single processor system, critical section problem can be solved by simply **disabling interrupt**
  - Non-preemptive kernel
- Problems
  - Inefficiency in multiprocessor environment
  - In some systems, clock is updated by interrupt
- This approach is taken by non-preemptive kernels.

# Hardware Instructions

공식에서 연산 한번에 실행

- If H/W provides **atomic instructions**, locking is easier to implement

- TestAndSet: set a variable to true and returns its previous value

```
boolean TestAndSet(boolean *target) {  
    boolean rv = *target;  
    *target = true;  
    return rv;  
}
```

(lock 변수) function x  
가(가)를 변경  
한번 실행하는 일까지 실행  
중간기 종료 x

- Swap: exchanges two variables

```
void Swap(boolean *a, boolean *b) {  
    boolean temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

lock 변수, local 변수

# Hardware Instructions

- If H/W provides **atomic instructions**, locking is easier to implement
  - Compare and swap (CAS): The operand value is set to new value only if the expression (`*value == expected`) is true.

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
    if (*value == expected)  
        *value = new_value;  
    return temp;  
}
```



# Mutual Exclusion using TestAndSet

- Shared variable
  - boolean lock = false;

## ■ Process $P_i$

do {

*Entry section* while (TestAndSet(&lock));

*return value == 0 (set)*  
*Atomic instruction*

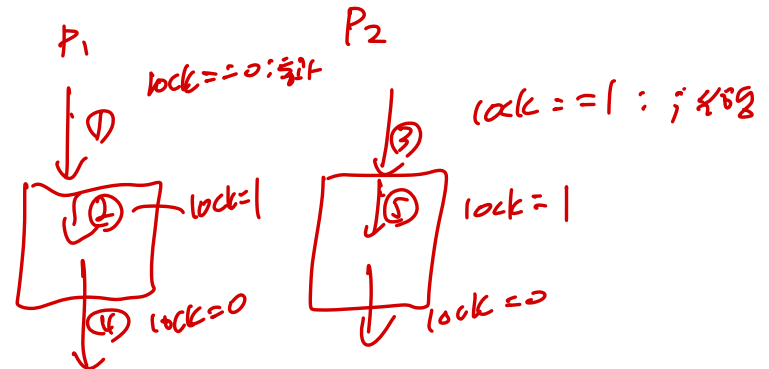
*/\* critical section \*/*

lock = false;

*/\* remainder section \*/*

} while(1);

*lock = 0 / 0*

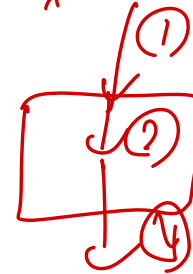


# Mutual Exclusion using Swap

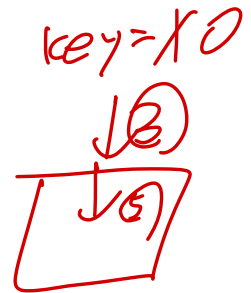
- Shared variable
  - boolean lock = false;

- Process  $P_i$ 
  - do {
    - key = true;
    - while (key == true)
      - Swap(&lock, &key);
    - /\* critical section \*/
    - lock = false;
    - /\* remainder section \*/
  - } while(1);

lock = false  
key = 10



// local var.



# Mutual Exclusion using CAS

- Shared variable

- boolean lock = false;

- Process  $P_i$

```
while (true) {
```

```
    while(compare_and_swap(&lock, 0, 1) != 0);
```

```
    /* critical section */
```

```
    lock = 0;
```

```
    /* remainder section */
```

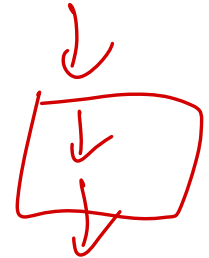
```
}
```

$lock = 0 \neq 1 \neq 0$

$P_1$



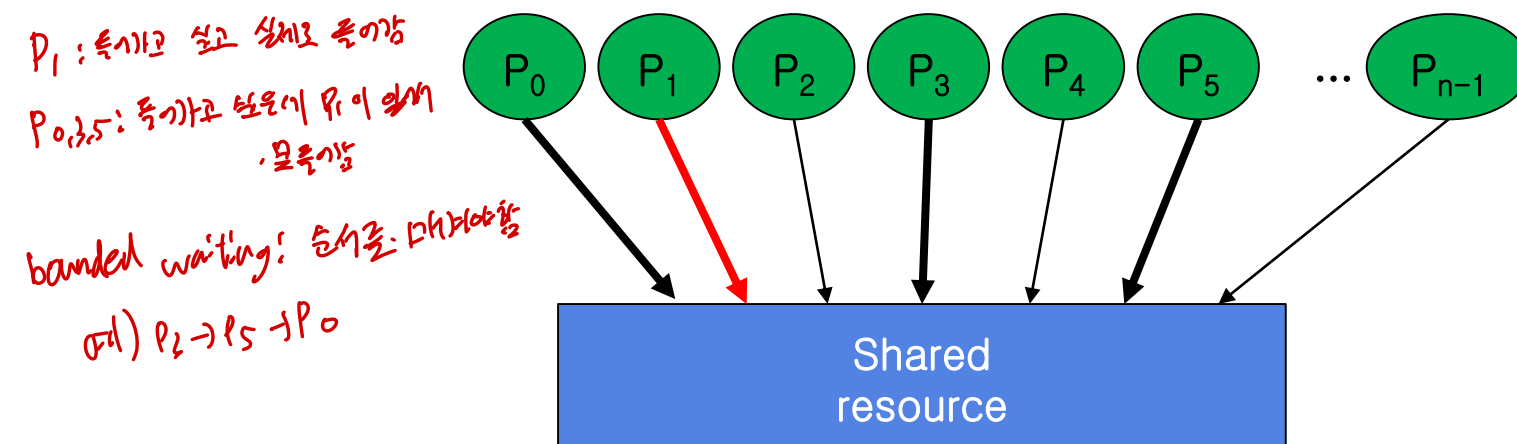
$P_2$



# Bounded Waiting Mutual Exclusion

- Bounded waiting for  $n$  processes

- Some processes want to enter their critical sections  
Ex)  $P_0, P_1, P_3, P_5$
- One of them, (ex:  $P_1$ ) is in its critical section



- Idea: Whenever a process leaves its critical section, it designates the next process to enter the critical section.
  - The next process: The first waiting process in right direction.

# Bounded Waiting Mutual Exclusion

각 변수가 어떤 목적으로 활용하

## Shared variables

- boolean lock;
  - For mutual exclusion
- boolean waiting[n];
  - if waiting[i] is true, it means  $P_i$  wants to enter critical section, but it didn't enter yet.

## Algorithm

lock = 0

```
while(TRUE){
    waiting[i] = TRUE;
    key = TRUE;      // local var.
    while(waiting[i] && key)
        key = TestAndSet(&lock);
    waiting[i] = FALSE;
```

entry section

/\* critical section \*/

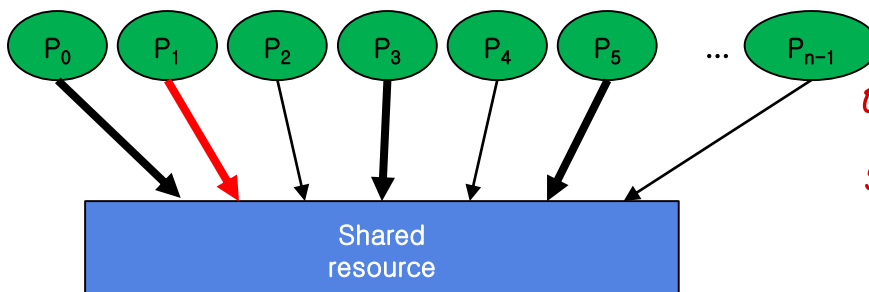
```
j = (i + 1) % n;
while(j != i && !waiting[j])
    j = (j + 1) % n;
```

```
if(j == i)
    lock = FALSE;
else
    waiting[j] = FALSE;
```

/\* remainder section \*/

}

exit section



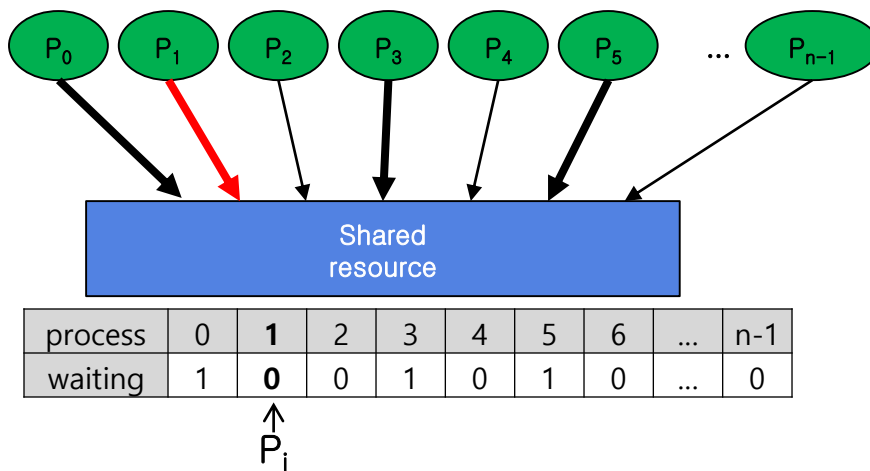
process	0	1	2	3	4	5	6	...	n-1
waiting	1	0	0	1	0	1	0	...	0

$\uparrow$   
 $P_i$

# Bounded Waiting Mutual Exclusion

## ■ Shared variables

- boolean lock;
  - For mutual exclusion
- boolean waiting[n];
  - if waiting[i] is true, it means  $P_i$  wants to enter critical section, but it didn't enter yet.



## ■ Algorithm using CAS

```
while(true) {
    waiting[i] = true;
    key = 1;
    while (waiting[i] && key)
        key =
            compare_and_swap(&lock,0,1);
    waiting[i] = false;
```

*/\* critical section \*/*

```
j = (i + 1) % n;
while ((j != i) && !waiting[j])
    j = (j + 1) % n;
if (j == i)
    lock = 0;
else
    waiting[j] = false;
```

*/\* remainder section \*/*

}

# Atomic Variables

- **Atomic variables**: variables that provide **atomic operations** on basic data types such as integers and Booleans

- Provided as “special data type + operation”
- Can be used in to ensure mutual exclusion

$V = 5$

- Typically, implemented by CAS instruction

```
void increment(atomic int *v)
```

```
{ increment 자르는 atomic 할당
```

```
    int temp = 0;
```

```
    do {
```

```
        temp = *v;
```

```
    } while (temp != compare_and_swap(v, temp, temp+1));
```

```
}
```

$P_1$   
temp=5  
 $V=5$   
 $V=6$

$P_2$   
temp=5  
 $V=6$   
 $P_1$ 이  $V=6$ 이므로  
temp=6  
 $V=6$   
 $V=7$

*atomic*

*$V$ 을 5로 바꾸고 6으로 올림*

# Agenda

---



- Background
- The critical-section problem
- H/W Support for Synchronization
- Mutex and Semaphore
- Monitors
- Liveness



# Mutex Locks

- A synchronization tool to implement mutual exclusion

- A Boolean variable *available* *available 하러 갔으면 lock, lock 변수의 반대*
  - Opposite to the lock variable in the previous section

- Atomic operations on mutex locks

- acquire() function *entry section의 시작*

```
acquire() {  
    while (!available);           /* busy waiting */  
    available = false;  
}
```
- release() function *exit section의 종료*

```
release() {  
    available = true;  
}
```

# Mutex Locks

- Solution to the critical-section problem using mutex lock

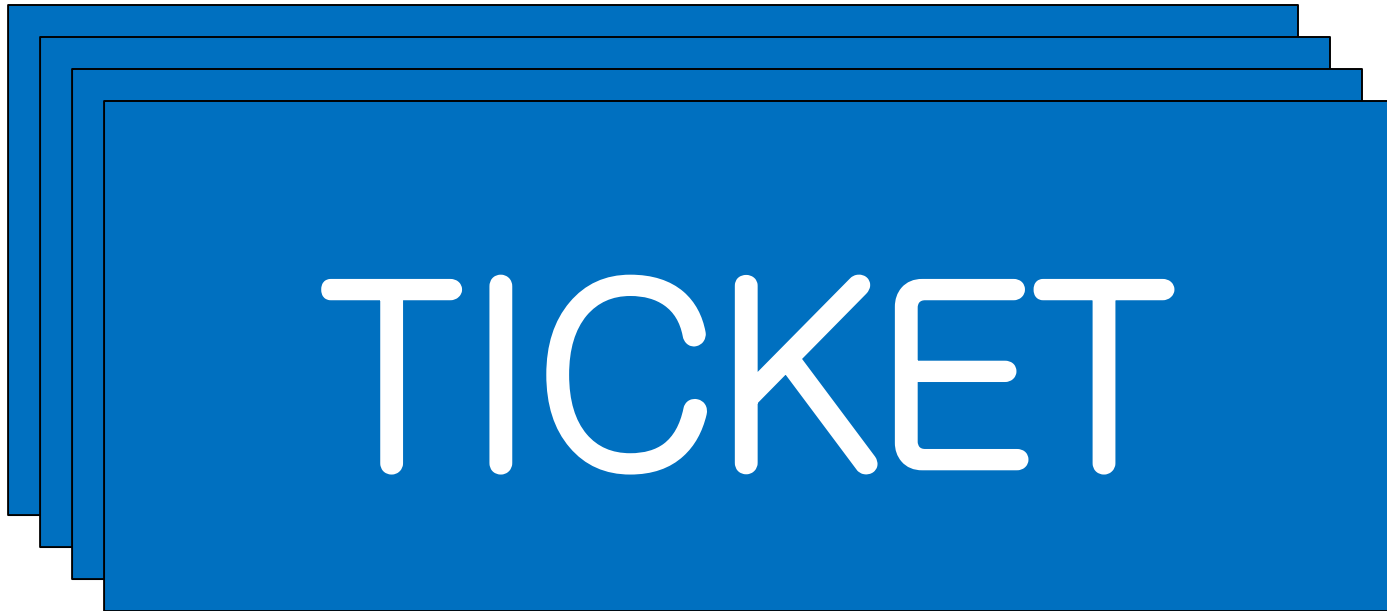
```
do {  
    mutex.acquire()  
    critical section  
    mutex.release()  
    remainder section  
} while(TRUE);
```

# Semaphores

- **Semaphore**: an integer variable accessed only through two **atomic operations**: **wait()** and **signal()**
  - Initial value of S is 1 or a positive integer.
  - Wait // correspond to mutex.acquire()  
wait(S) {  
    while(S <= 0); // waits for the lock  
    S--; // holds the lock  
}
  - Signal // correspond to mutex.release()  
signal(S) {  
    S++; // releases the lock  
}

# Semaphores

- S is similar to the number of reusable tickets to enter the critical section.



# Semaphores

## ■ Synchronization using semaphore

```
do {  
    wait(S);  
    critical section  
    signal(S);  
    remainder section  
} while(1);
```

- If initially value of S is 1,
  - wait(S) acquires lock
  - signal(S) releases lock

```
■ Wait  
wait(S)  
{  
    while(S <= 0);  
    S--;  
}
```

```
■ Signal  
signal(S)  
{  
    S++;  
}
```

# Implementation of Semaphore

- Problem of previous definition of wait(): spinlock

```
wait(S) {  
    while(S <= 0);           // spinlock (busy waiting)  
    S--;  
}
```

- Alternative implementation: **block** instead of spinlock
  - If a process invokes wait(S), put the process into a waiting queue of S and block itself.

# Implementation of Semaphore



## ■ New definition of semaphore

```
typedef struct {  
    int value;
```

```
    struct process *list;    // head of linked list for waiting queue  
} semaphore;
```

*PCB*      *ready queue → linked list 구조*

## ■ Wait

```
wait(semaphore *S){  
    S->value--;  
    if(S->value < 0){  
        add this process to S->list;  
        block(); // suspend  
    }  
}
```

## ■ Signal

```
signal(semaphore *S)  
{  
    S->value++;  
    if(S->value <= 0){  
        remove a process P from S->list;  
        wakeup(P); // resume  
    }  
}
```

# Implementation of Semaphore



- A critical aspect of semaphore: **atomicity**
  - Atomicity is often enforced by mutual exclusion → *atomic한 것과 같은 효과*
  - Single processor environment: disable interrupt → *mutual exclusion* → *atomic*
  - Multiprocessor environment: spinlock
    - Disable interrupt of all processors → performance degradation
    - Spinlock is much shorter than previous algorithms



# Agenda

---



- Background
- The critical-section problem
- H/W Support for Synchronization
- Mutex and Semaphores
- Monitors
- Liveness

# Monitors

- Motivation: semaphore is still too low-level tool.

Example)

correct

```
wait(mutex);  
...  
    critical section  
...  
signal(mutex);
```

wrong

```
signal(mutex);  
...  
    critical section  
...  
wait(mutex)
```

wrong

```
wait(mutex);  
...  
    critical section  
...  
wait(mutex);
```

- **Monitor**: a high-level language construct to support synchronization
  - Private data: accessible only through the methods
  - Public methods: mutual exclusion is provided

# Class/Structures in OOP Language

## ■ C language

```
struct Stack {  
    int array[100];  
    int top;  
};  
  
void Push(struct Stack *s,  
    int item);  
int Pop(struct Stack *s);  
  
struct Stack stack1;  
Push(&stack1, 10);
```

## ■ C++, Java

```
class Stack {  
    int array[100];  
    int top;  
private:  
    void Push(int item);  
    int Pop();  
};  
  
Stack stack2;  
stack2.Push(10);
```

# Monitors

## ■ Syntax of monitor

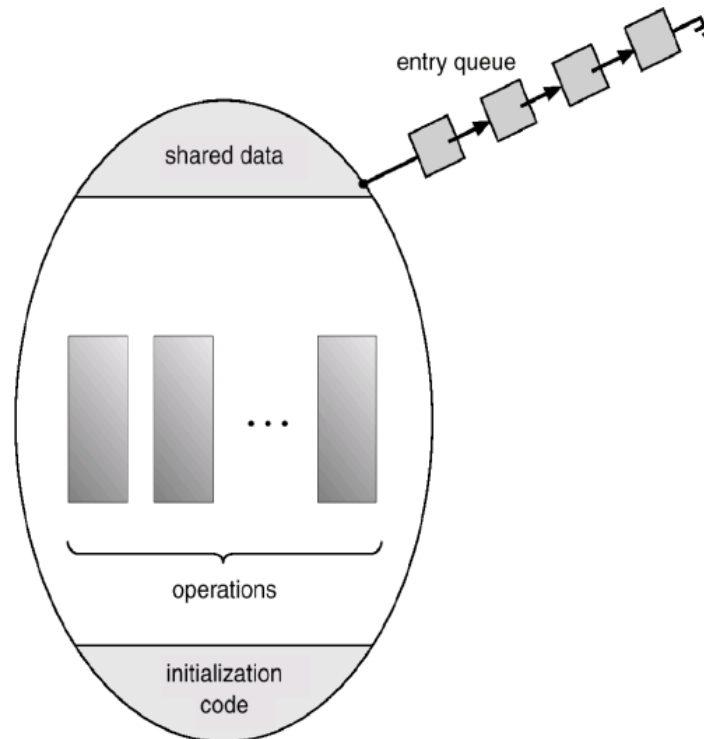
```
monitor monitor-name
{
    // shared variable declarations

    procedure body  $P1$  (...) {
        ...
    }
    procedure body  $P2$  (...) {
        ...
    }
    ...
    procedure body  $Pn$  (...) {
        ...
    }
    initialization code (...) {
        ...
    }
}
```

Only one process can be active  
within the monitor at a time

# Monitors

- Only one process can be active within the monitor at a time.
  - The programmer doesn't have to concern about synchronization.



# Monitor in Java

## ■ Synchronized method

Ex)

```
class Producer {  
    private int product;  
    ...  
    private synchronized void produce();    // mutually exclusive  
}
```

mutex semaphore 같은 것

c.f. C#: System.Threading.Monitor class

# Condition Variables

## ■ Condition: additional synchronization mechanism

- Variable declaration

condition x, y;

- Wait

x.wait();

// invoking process is suspended until

// other process call x.signal()

- Signal

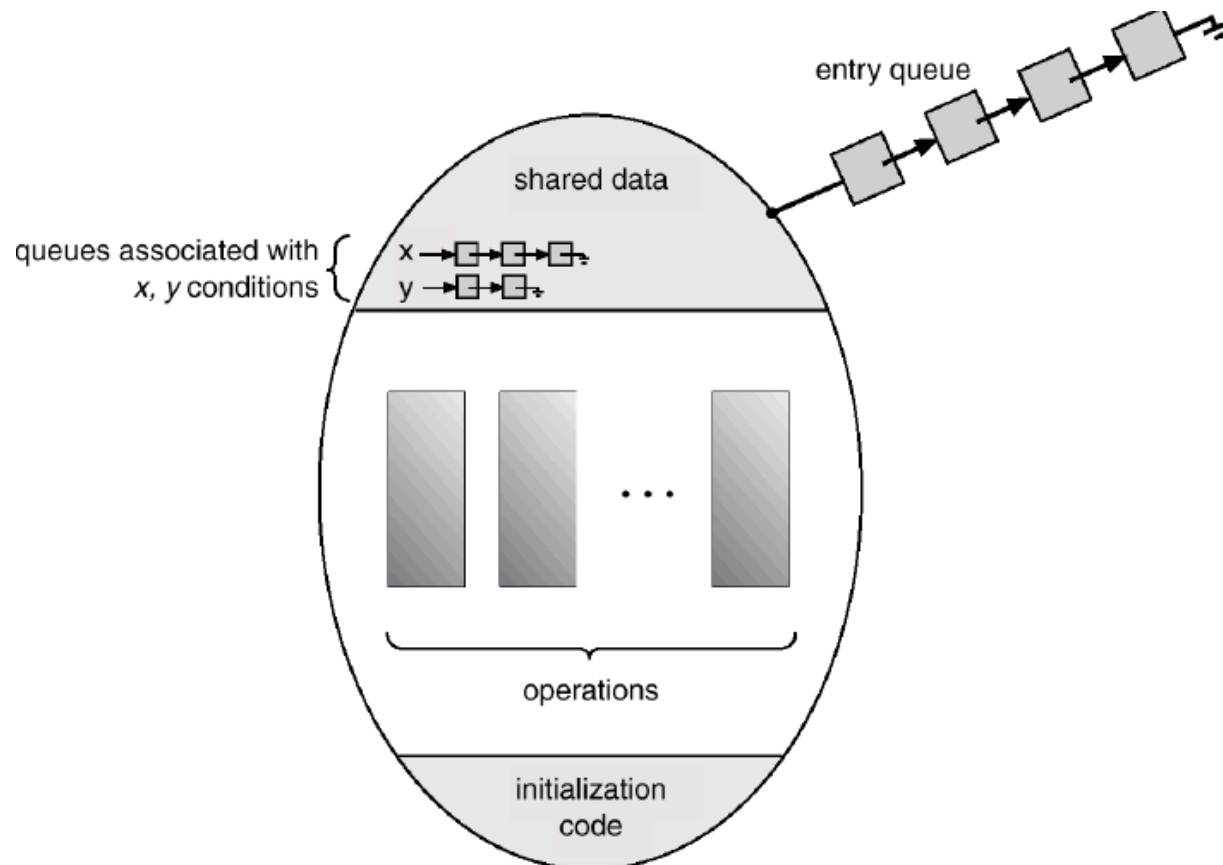
x.signal();

// resumes exactly one suspended process

// if no process is waiting, do nothing

*Semaphores 2302*

# Condition Variables





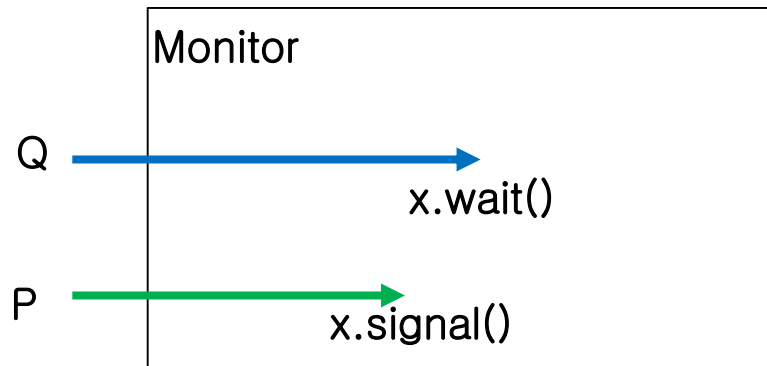
# Condition Variables

## ■ Problem

- Process P wakes up another process Q by invoking `x.signal()`.
- Both P and Q are executing in monitor.

## ■ Solution

- **Signal and wait**: P waits until Q leaves monitor or waits
- **Signal and continue**: Q waits until P leaves monitor or waits



# Implementing a Monitor Using Semaphores

---



- Functions to implement
  - External procedures
  - wait() of condition variable
  - signal() of condition variable
  
- Two semaphores are required:
  - semaphore *mutex* = 1; // mutual exclusion
  - semaphore *next* = 0; // signaling process should wait  
// until the resumed process leaves monitor
  - int next\_count = 0;

# Implementing a Monitor Using Semaphores

---



## ■ External procedure $F$

```
wait(mutex);
```

```
...
```

```
body of  $F$ ;
```

```
...
```

```
signal(mutex);
```

# Implementing a Monitor Using Semaphores



## ■ Condition variables

### ■ Required data

- semaphore x\_sem; // (initially, 0)
- int x\_count = 0;

## ■ x.wait()

```
{
    x_count++;
    if (next_count > 0)
        signal(next);
    else
        signal(mutex);
    wait(x_sem);
    x_count--;
}
```

## ■ x.signal()

```
{
    if (x_count > 0) {
        next_count++;
        signal(x_sem);
        wait(next);
        next_count--;
    }
}
```

# Resuming Processes within a Monitor

- If a process invoke `x.signal()` and there are several processes waiting on `x`, which one should be resumed?
  - Sometimes, FIFO is not enough
- Conditional-wait
  - `x.wait(c);`            `// c: priority number`

# Agenda

---



- Background
- The critical-section problem
- H/W Support for Synchronization
- Mutex and Semaphores
- Monitors
- Liveness

# Liveness



---

- **Liveness**: a set of properties that a system must satisfy to ensure that processes **make progress** during their execution life cycle.
- Reasons of liveness failure
  - Deadlock
  - Priority inversion
  - Many other reasons

# Deadlock and Starvation

## ■ Deadlock

$P_0$	$P_1$
<i>wait(S);</i>	<i>wait(Q);</i>
<i>wait(Q);</i>	<i>wait(S);</i>
$\vdots$	$\vdots$
<i>signal(S);</i>	<i>signal(Q);</i>
<i>signal(Q);</i>	<i>signal(S);</i>



## ■ Starvation (infinite blocking): a situation in which a process waits indefinitely within the semaphore

Ex) waiting list is implemented by LIFO order



# Priority Inversion



- **Priority inversion**: a higher-priority process needs to read or modify kernel data that are currently being accessed by a (chain of) lower-priority process
  - Three processes L, M, and H
    - Priorities:  $L < M < H$
  - H waits for a semaphore S acquired by L.
  - L should wait for M because of its lower priority
  - Consequently, H waits for M
- Can be avoided by priority-inheritance protocol