

# 7. Synchronization Examples

ECE30021/ITP30002 Operating Systems

# Agenda

---



- Classical Problems of Synchronization
- Synchronization within the Kernel
- POSIX Synchronization
- Synchronization in Java
- Alternative Approaches

# Classical Problems of Synchronization

---



- The bounded-buffer problem
- The readers-writers problem
- The dining-philosophers problem

# The Bounded-Buffer Problem



- If the buffer is full, the producer must wait until the consumer deletes an item.
  - Producer needs an empty space
  - **# of empty slot** is represented by a semaphore *empty*
- If the buffer is empty, the consumer must wait until the producer adds an item.
  - Consumer needs an item
  - **# of item** is represented by a semaphore *full*

# The Bounded-Buffer Problem

- Producer-consumer problem with bounded buffer

- Semaphores:  $\text{full} = 0$ ,  $\text{empty} = n$ ,  $\text{mutex} = 1$ ;

- Producer


```
do {  
    ...  
    produce an item in nextp  
    ...  
    wait(empty);  
    wait(mutex);    // acquire()  
    ...  
    add nextp to buffer  
    ...  
    signal(mutex);  // release()  
    signal(full);  
} while (1);
```

- Consumer

```
do {  
    wait(full);  
    wait(mutex);    // acquire()  
    ...  
    remove an item from buffer to nextc  
    ...  
    signal(mutex);  // release()  
    signal(empty);  
    ...  
    consume the item in nextc  
    ...  
} while (1);
```

# The Readers–Writers Problem

- There are multiple readers and writers to access a shared data
  - Readers can access database simultaneously.
  - When a writer is accessing the shared data, no other thread can access it.



# The Readers–Writers Problem



## ■ Behavior of a writer

- If a thread is in the critical section, all writers must wait.
- The writer can enter the critical section only when no thread is in its critical section.
  - It should prevent all threads from entering the critical section.

## ■ Behavior of a reader

- If no writer is in its critical section, the reader can enter the critical section.
- Otherwise, the reader should wait until the writer leaves the critical section.
- When a reader is in its critical section, any reader can enter the critical section, but no writer can.
  - Condition for the first reader is different from the following readers.

# The Readers–Writers Problem

## ■ Shared data

- semaphore mutex=1, wrt=1;
- int readcount = 0;
  - # of readers in critical section

## ■ Writer

wait(wrt);

...

writing is performed

...

signal(wrt);

## ■ Reader

wait(mutex);

readcount++;

if (readcount == 1)

wait(wrt);

signal(mutex);

...

reading is performed

...

wait(mutex);

readcount--;

if (readcount == 0)

signal(wrt);


signal(mutex);



# The Dining Philosophers Problem

## ■ Problem definition

- 5 philosophers sitting on a circular table
  - Thinking or eating
- 5 bowls, 5 single chopsticks
- No interaction with colleagues
- To eat, the philosopher should pick up two chopsticks closest to her
- A philosopher can pick up only one chopstick at a time
- When she finish eating, she release chopsticks



- Solution should be deadlock-free and starvation free

# The Dining Philosophers Problem

- A possible solution, but **deadlock can occur**

```
do {  
    wait(chopstick[i]);           // pick up left chopstick  
    wait(chopstick[(i+1) % 5]);   // pick up right chopstick  
    ...  
    eat  
    ...  
    signal(chopstick[i]);         // release left chopstick  
    signal(chopstick[(i+1) % 5]); // release right chopstick  
    ...  
    think  
    ...  
} while (TRUE);
```

# Deadlock

## ■ Deadlock

$P_0$	$P_1$
$wait(S);$	$wait(Q);$
$wait(Q);$	$wait(S);$
$\vdots$	$\vdots$
$signal(S);$	$signal(Q);$
$signal(Q);$	$signal(S);$



## ■ Necessary conditions for deadlock

1. Mutual exclusion
2. Hold and wait *물건이 2개는 잡는다*
3. No preemption *임의로 빼앗아가지 않음*
4. Circular wait *각각의 자원을, 다른 자원을*

# Dining Philosophers Solution Using Monitor

---



## ■ Data structures

- `enum { thinking, hungry, eating } state[5];`
- `condition self[5];`

## ■ Process of i-th philosopher implemented by a monitor dp

```
dp.pickup(i);           // entry section
...
Eat                     // critical section
...
dp.putdown(i);          // exit section
```

# Dining Philosophers Solution Using Monitor

```
monitor diningPhilosophers {  
    int state[5];  
    static final int THINKING = 0;  
    static final int HUNGRY = 1;  
    static final int EATING = 2;  
    condition self[5];
```


```
    void initialization_code {  
        for (int i = 0; i < 5; i++)  
            state[i] = THINKING;  
    }
```

```
    void pickUp(int i) {  
        state[i] = HUNGRY;  
        test(i);  
        if (state[i] != EATING)  
            self[i].wait();  
    }
```

```
    void putDown(int i) {  
        state[i] = THINKING;  
        // test left and right  
        test((i+4) % 5);  
        test((i+1) % 5);  
    }
```

```
    void test(int i) {  
        if((state[(i+4) % 5] != EATING) &&  
            (state[i] == HUNGRY) &&  
            (state[(i+1) % 5] != EATING) ) {  
            state[i] = EATING;  
            self[i].signal();  
        }  
    }
```

Deadlock-free, but not starvation-free



# Agenda

---



- Classical Problems of Synchronization
- Synchronization within the Kernel
- POSIX Synchronization
- Synchronization in Java
- Alternative Approaches


# Synchronization in Windows Kernel



- Windows kernel
  - A multithreaded kernel
  - Provides support for real-time applications
  - Provides support for multiple processors
  
- Accessing global variables
  - Single-processor system: mask interrupt (Kernel mode interrupt disable)
  - Multi-processor system: spinlocks
    - Only to protect short code segments
    - The kernel ensures that a thread will never be preempted while holding a spinlock

# Synchronization in Windows

- For thread synchronization **outside the kernel**, windows provides **dispatcher objects**
  - **Mutex locks, semaphores, events, timers**
- **States of dispatcher objects**
  - **Signaled state**: the object is available
    - Thread will not block when acquiring the object.
  - **Nonsignaled state**: the object is **not** available
    - Thread will block when attempting to acquire the object.
    - Its state changes from ready to waiting, and the thread is placed in a waiting queue for that object.





# Synchronization in Windows



- **Critical section object:** a user-mode mutex that can often be acquired and released without kernel intervention.
  - A critical-section object **first uses a spinlock** while waiting for the other thread to release the object.
    - ➔ **Does not allocate kernel object. (efficient)**
  - If it spins too long, the acquiring thread will then allocate a kernel mutex and yield its CPU.

# Synchronization in Linux Kernel

- From v.2.6., Linux kernel is fully preemptive
- Atomic integer (atomic\_t)
  - All math operations are performed without interruption

Ex) atomic\_t counter;  
int value;

<i>Atomic Operation</i>	<i>Effect</i>
<code>atomic_set(&amp;counter,5);</code>	<code>counter = 5</code>
<code>atomic_add(10,&amp;counter);</code>	<code>counter = counter + 10</code>
<code>atomic_sub(4,&amp;counter);</code>	<code>counter = counter - 4</code>
<code>atomic_inc(&amp;counter);</code>	<code>counter = counter + 1</code>
<code>value = atomic_read(&amp;counter);</code>	<code>value = 12</code>

# Synchronization in Linux Kernel



## ■ Mutex locks

- `int mutex_init(mutex_t *mp, int type, void *arg);`
- `int mutex_lock(mutex_t *mp);`
- `int mutex_trylock(mutex_t *mp);`
- `int mutex_unlock(mutex_t *mp);`
- `int mutex_consistent(mutex_t *mp);`
  - Makes a mutex consistent
- `int mutex_destroy(mutex_t *mp);`

# Synchronization in Linux Kernel



## ■ Spinlocks

- `spin_lock()`, `spin_unlock()`, etc.
- Useful for short duration
- Inappropriate on single processing core

## ■ Enabling/disabling kernel preemption

- `preempt_disable()`, `preempt_enable()`

Single processor	Multiprocessor
Disable kernel preemption	Acquire spinlock
Enable kernel preemption	Release spinlock

# Agenda

---



- Classical Problems of Synchronization
- Synchronization within the Kernel
- POSIX Synchronization
- Synchronization in Java
- Alternative Approaches

# POSIX Synchronization

## ■ POSIX mutex locks

### ■ Creation and initialization

#include <pthread.h>

```
pthread_mutex_t mutex;           // global declaration  
pthread_mutex_init(&mutex, NULL); // call before first lock
```

*pthread\_mutex\_t;*  
or *pthread\_mutex\_destroy();*

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

### ■ Critical section

```
pthread_mutex_lock(&mutex);
```

```
/* critical section */
```

```
pthread_mutex_unlock(&mutex);
```

# POSIX Synchronization

## ■ Named semaphores

- Multiple unrelated processes can easily use a common semaphore

- Creation and initialization

```
#include <semaphore.h>
```

```
sem_t *sem; // global declaration
```

```
sem = sem_open("SEM", O_CREAT, 0666, 1);
```

- Critical-section

```
wait sem_wait(sem); // acquire the semaphore
```

```
/* critical section */
```

```
signal sem_post(sem); // release the semaphore
```

```
named semaphore: sem_close(), sem_unlink()
```

# POSIX Synchronization

## ■ Unnamed semaphores

- Creation and initialization

```
#include <semaphore.h>
```

```
sem_t sem;           // global declaration
```

```
sem_init(&sem, 0, 1);  
            ?
```

unnamed semaphore:  
sem\_destroy()

- Critical-section

```
sem_wait(&sem);       // acquire the semaphore
```

```
/* critical section */
```

```
sem_post(&sem);       // release the semaphore
```



# POSIX Synchronization



## ■ Condition variables

- Combined with mutex locks instead of monitors
- Initialization

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;
```

```
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```

- Critical-section

### Thread A

```
pthread_mutex_lock(&mutex);  
while (a != b)  
    pthread_cond_wait(&cond_var, &mutex);  
pthread_mutex_unlock(&mutex);
```

### Thread B

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```

# Agenda

---




- Classical Problems of Synchronization
- Synchronization within the Kernel
- POSIX Synchronization
- Synchronization in Java
- Alternative Approaches

# Synchronization in Java

## ■ Java monitors

- Synchronized methods
- Entry set
  - Threads waiting for the lock
- Releasing the lock
  - Resumes an **arbitrarily** chosen thread
  - In practice, FIFO policy



```
public class BoundedBuffer<E>
{
    private static final int BUFFER_SIZE = 5;

    private int count, in, out;
    private E[] buffer;

    public BoundedBuffer() {
        count = 0;
        in = 0;
        out = 0;
        buffer = (E[]) new Object[BUFFER_SIZE];
    }

    /* Producers call this method */
    public synchronized void insert(E item) {
        /* See Figure 7.11 */
    }

    /* Consumers call this method */
    public synchronized E remove() {
        /* See Figure 7.11 */
    }
}
```

# Synchronization in Java

## ■ Producer-Consumer in Java

```
/* Producers call this method */
public synchronized void insert(E item) {
    while (count == BUFFER_SIZE) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
    count++;

    notify();
}
```

```
/* Consumers call this method */
public synchronized E remove() {
    E item;

    while (count == 0) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;

    notify();

    return item;
}
```

# Synchronization in Java



## ■ Java monitors

- If the thread that has the lock is unable to continue, it calls `wait()`.
  - Ex) Calling `insert()` when the buffer is full
  - 1. The thread releases the lock for the object.
  - 2. The state of the thread is set to blocked.
  - 3. The thread is placed in the `wait set` for the object.
- Other thread can call `notify()` to wake up a waiting thread
  - 1. Picks an arbitrary thread T from the list of threads in the `wait set`
  - 2. Moves T from the `wait set` to the `entry set`
    - When the lock is released, JVM arbitrarily chooses a thread to run
  - 3. Sets the state of T from blocked to runnable

# Synchronization in Java

## ■ Semaphores

- Constructor
  - Semaphore(int value);
- Critical-section

```
Semaphore sem = new Semaphore(1);

try {
    sem.acquire();
    /* critical section */
}
catch (InterruptedException ie) { }
finally {
    sem.release();
}
```

# Agenda

---



- Classical Problems of Synchronization
- Synchronization within the Kernel
- POSIX Synchronization
- Synchronization in Java
- Alternative Approaches

# Transactional Memory

- **Memory transaction**: a sequence of memory read–write operations that are atomic.
  - Originated from database theory
  - SW transactional memory(STM) / HW transactional memory(HTM)

Ex)

```
void update ()
{
    atomic {
        /* modify shared data */
    }
}
```

- Traditional implementation

```
void update ()
{
    acquire();
    /* modify shared data */
    release();
}
```



# OpenMP



## ■ Parallel region

- #pragma omp parallel

## ■ Critical section

- #pragma omp critical

Ex)

```
void update(int value)
{
    #pragma omp critical
    {
        counter += value;
    }
}
```