

# 10. Virtual Memory

ECE30021/ITP30002 Operating Systems

# Agenda

---



- Background
- Demand paging
- Copy-on-Write (COW)
- Page replacement
- Allocation of frames
- Thrashing
- Memory-mapped files
- Allocating kernel memory
- Other considerations
- Operation-system examples

# Background



- Instructions should be in physical memory to be executed
  - ➔ In order to execute a program, should we load **entire program** in memory?
- Some parts are rarely used
  - Error handling codes
  - Arrays/lists larger than necessary
  - Rarely used routines
- Alternative
  - Executing program which is only partially in memory

# Background


---



- If we can run a program by loading in parts ...
  - A program is not constrained by the amount of physical memory
  - More program can run at the same time
  - Less I/O is need to load or swap programs

# Virtual Memory

- **Virtual memory:** a separation of logical memory from physical memory
  - Contents not currently reside in main memory can be addressed.  
➔ H/W and OS will load the required memory from auxiliary storage automatically.
  - User programs can reference more memory than actually exists



# Virtual Memory




- **Virtual memory**: a separation of logical memory from physical memory
  - Only part of the program needs to be in memory for execution
  - Logical address space can be much larger than physical address space
  - Allows address spaces to be shared by several processes
  - Allows for more efficient process creation (e.g. COW)
  - More programs running concurrently
  - Less I/O needed to load or swap processes

*page 251*


# Virtual Address Space

- **Virtual address space:** logical view of how process is stored in memory
  - Usually start at address 0, contiguous addresses until end of space.
  - Meanwhile, physical memory organized in page frames
  - MMU must map logical to physical
  - Programmers don't have to concern about memory management
  - Virtual address space can be sparse



# Shared Library Using Virtual Memory

- Shared page






# Implementation Techniques

## ■ Demand paging

- Paging + swapping

## ■ Demand segmentation

- Segmentation + swapping



# Agenda

---




- Background
- Demand paging
- Copy-on-Write (COW)
- Page replacement
- Allocation of frames
- Thrashing
- Memory-mapped files
- Allocating kernel memory
- Other considerations
- Operation-system examples

# Demand Paging

- Similar to paging system with swapping


- Lazy swapper (or pager)  
*꼭 필요할때만 하거나* *page swapper*

- Pages are only loaded when they are demanded during execution.



# Demand Paging

- Pages are loaded **only when they are demanded** during program execution
  - Requires H/W support to distinguish the page on memory or on disk



# Page Table with Valid/Invalid Bit

- Valid/invalid bit of each page
  - Valid: the page is valid and exists in physical memory
  - Invalid: the page is not valid (not in the valid logical address space of the process) or not loaded in physical memory
- If program tries to access ..
  - Valid page: execution proceeds normally
  - Invalid page: cause **page-fault** trap to OS

Frame #	valid-invalid bit
	v
	v
	v
	i
...	
	i
	i

page table


# Basic Concepts

page fault 발생 | valid-invalid bit 0 | 잘못 가져옴. [재실행]  
3EX

## ■ Handling page-fault

- Check an internal table to determine whether the reference was valid or not
- If the reference was invalid, terminates the process
- If valid, page it in.
  - Find a free frame
  - Read desired page into the free frame
  - Modify internal table
  - Restart the instruction that caused the page-fault trap

# Handling Page-Fault



# Basic Concepts



- In pure demand paging, a page is not brought into memory until it is required
- Performance degradation
  - Theoretically, some program can cause multiple page faults per instructions
  - But actually, this behavior is exceedingly unlikely.
    - **locality of references**
- H/W supports for demand paging
  - Page table (support for valid-invalid bit, ...)
  - Secondary memory (swap space)
  - Instruction restart



# Performance of Demand Paging

## ■ Effective access time

Effective access time =  $(1-p) * ma + p * \text{<page fault time>}$

□  $ma$ : <sup>main</sup>memory access time (10~200 nano sec.)

□  $p$ : probability of page fault ~~ex)~~ 0.00

## ■ Page fault time

- Service page-fault interrupt → 1~100 μsec.
- Read in the page → about 8 msec
- Restart the process → 1~100 μsec.

# Performance of Demand Paging

Effective access time =  $(1-p) * ma + p * \text{<page fault time>}$

## ■ Example

- Memory access time: 200 nano sec
- Page-fault service time: 8 milliseconds

## ■ Then...

Effective access time (in nano sec.)

$$= (1-p) * 200 + 8,000,000 * p$$

$$\approx 200 + 7,999,800 * p$$

- Proportional to **page fault rate**

*ma x 1000*

Ex)  $p == 1/1000$ , effective access time = 8.2  $\mu$ sec. (40 times)

- Page fault rate should be kept low

# Execution of Program in File System



- Ways to execute a program in file system
  - Option1: copy entire file into swap space at starting time
    - Usually swap space is faster than file system
  - Option2: initially, demand pages from files system and all subsequent paging can be done from swap space
    - Only needed pages are read from file system

# Agenda

---



- Background
- Demand paging
- Copy-on-Write (COW)
- Page replacement
- Allocation of frames
- Thrashing
- Memory-mapped files
- Allocating kernel memory
- Other considerations
- Operation-system examples

# Copy-on-Write

## ■ fork() copies process

- Duplicates pages belong to the parent

process creation, DRAM 복사 overhead 최소화  
parent process가 child process로 복사되었을 때 논리적으로는 별다른 문제가 없으나 실제로는 같은 메모리 지소를 가진  
즉, 한 process가 writing을 하면 이 메모리 지소를 공유하고 있는 모든 process가 영향을 받는다.

## ■ Copy-on-write (COW)

- When the process is created, pages are not actually duplicated but just shared.
  - Process creation time is reduced.
- When either process writes to a shared page, a copy of the page is created.

cf. vfork() – logically shares memory with parent (obsolete)


## ■ Many OS's provides a list of free frames for COW or stack/heap that can be expanded

### ➔ Zero-fill-on-demand (ZFOD)

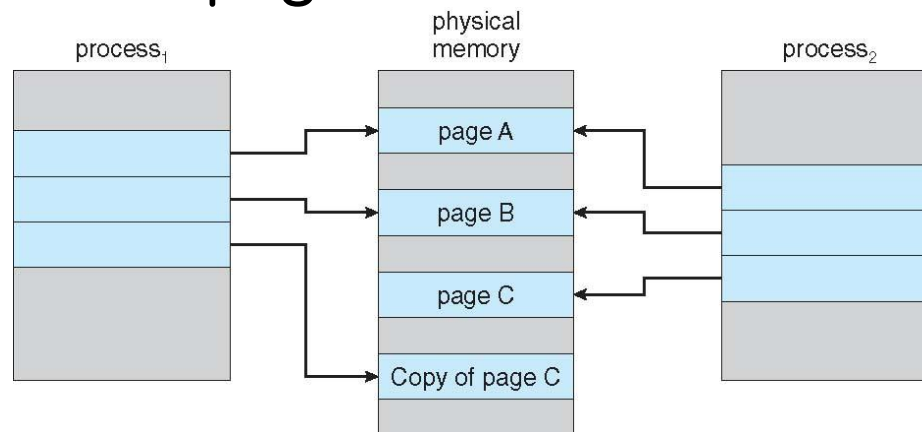
- Zero-out pages before being assigned to a process

# Copy-on-Write

- Before P1 modifies page C



- After P1 modifies page C



# Agenda

---



- Background
- Demand paging *paging + swapping*
- Copy-on-Write (COW)
- Page replacement
- Allocation of frames
- Thrashing
- Memory-mapped files
- Allocating kernel memory
- Other considerations
- Operation-system examples

# Two major problems



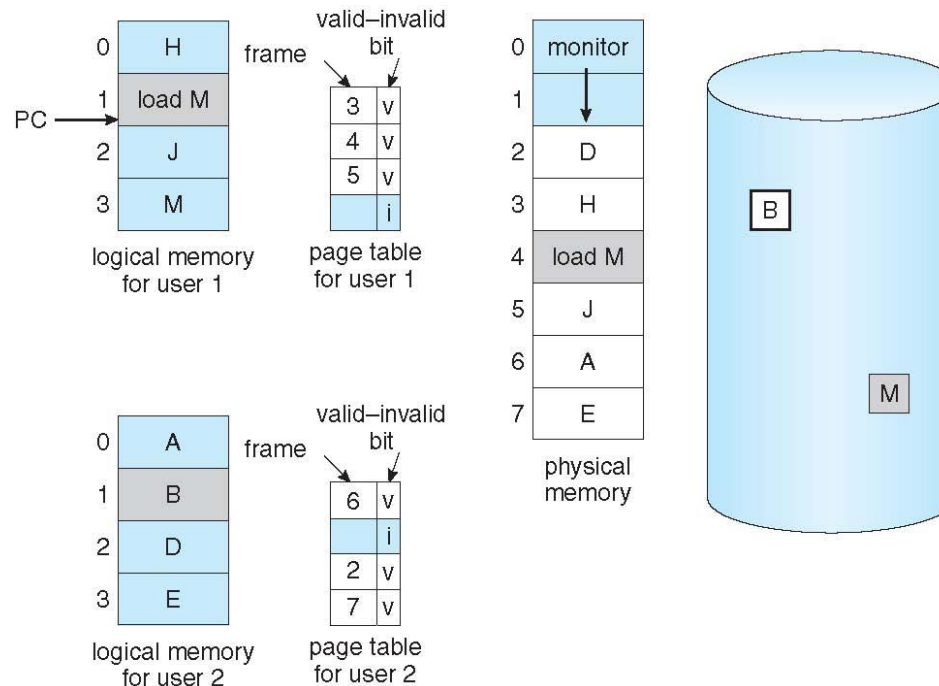
- Two major problems in demand paging
  - Page-replacement algorithm
  - Frame-allocation algorithms
- Even slight improvement can yield large gain in performance.
  - Effective access time =  $(1-p) * ma + p * \text{<page fault time>}$




# Page Replacement

## ■ Page replacement

- If no frame is free at a page fault, we find a frame not being used currently, and swap out
- Writing overhead can be reduced by **modify-bit** (or **dirty-bit**) for each frame




# Page Replacement



# Page Faults vs. Number of Frames

- In general, the more frames, the fewer page faults



# Page Replacement Algorithms



- FIFO page replacement
- Optimal page replacement (in theory)
- Least-recently-used (LRU) page replacement
- LRU-approximation page replacement
- ETC.

# FIFO Page Replacement

- **First-in, first-out:** when a page should be replaced, the oldest page is chosen.
  - Easy, but not always good

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2
	0	0	0
		1	1

2	2	4	4	4	0
3	3	3	2	2	2
1	0	0	0	3	3

0	0
1	1
3	2

7	7	7
1	0	0
2	2	1

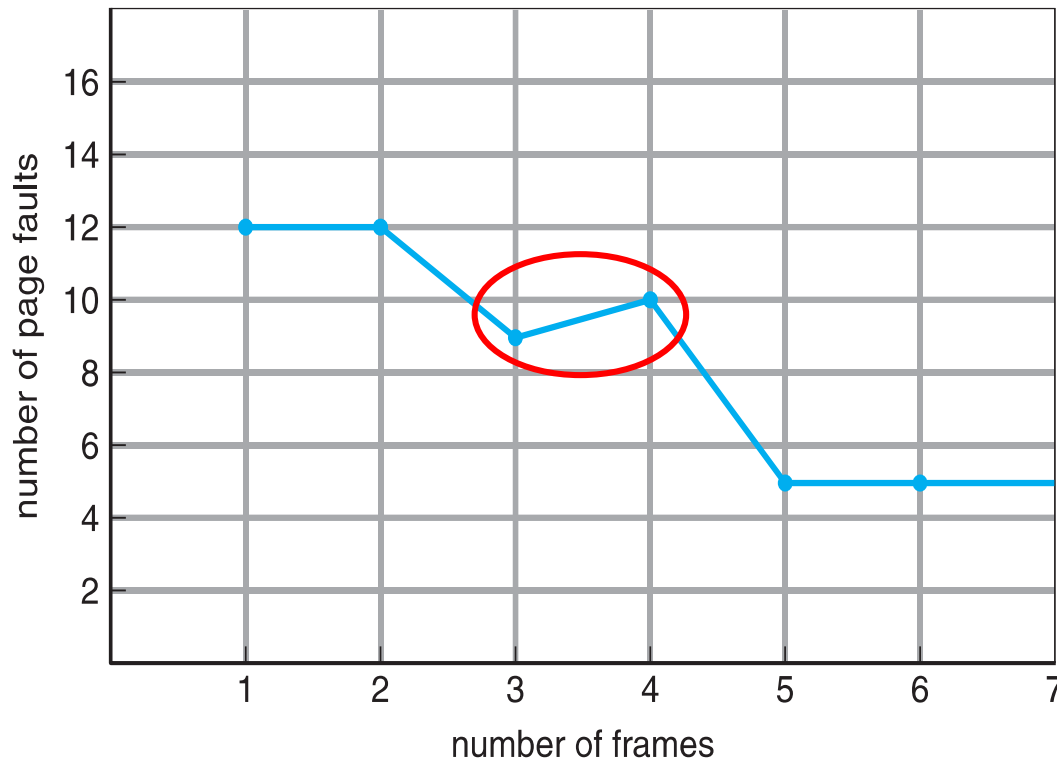
page frames

- # of page faults: 15
- **Problem: Belady's anomaly**

# FIFO Page Replacement

- **Belady's anomaly:** # of faults for 4 frames is greater than # of faults for 3 frames

(Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5)




Page-fault rate may increase as the number of frames increase.

# Optimal Page Replacement

- Replace the page that will not be used for the longest period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

- # of page faults: 9
- Problem: It requires future knowledge

# LRU Page Replacement

- LRU (Least Recently Used): replace page that has not been used for longest period of time



- # of page faults: 12
- LRU is considered to be good and used frequently




# Implementation of LRU

## ■ Using counter (logical clock)

- Associate with each page-table entry a **time-of-used field**
- Whenever a page is referenced, clock register is copied to its time-of-used field

## ■ Using stack of page numbers

- If a page is referenced, remove it and put on the top of the stack



6/25 느낌

↳ 항상 제일 최근에 참조된

# Stack Algorithm

- Does LRU cause the Belady's anomaly?
- Stack algorithm: an algorithm for which the set of pages in memory for  $n$  frames is always a subset of the set of pages that would be in memory with  $n + 1$  frames.
  - Never exhibit Belady's anomaly
  - LRU is a stack algorithm

$n$  frames

1	2	3	4	5	6	7	...

$n + 1$  frames

1	2	3	4	5	6	7	...

# LRU–Approximation Page Replacement

리스트 1번씩의 번호

## ■ Motivation

- LRU algorithm is good, but few system provide sufficient supports for LRU
- However, many systems support **reference bit** for each page
  - We can determine which pages have been referenced, but not their order.

## ■ LRU–approximation algorithms


- Additional–reference–bit algorithm
- Second–chance algorithm

# LRU–Approximation Page Replacement

## ■ Additional–reference–bit algorithm

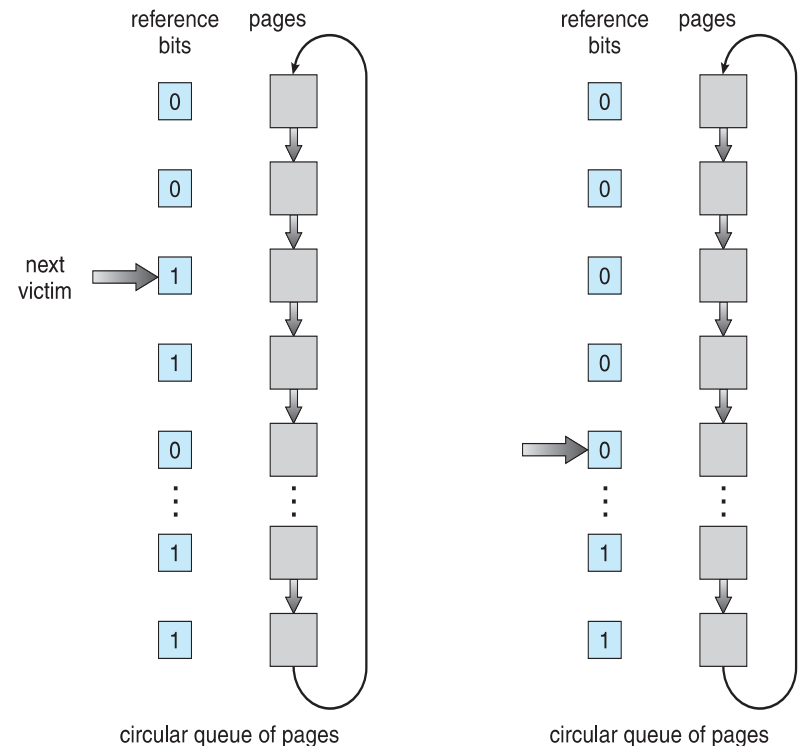
- Gain additional ordering information by recording the reference bits at regular intervals.

Ex) keep 8–bit byte for each page and records current state of reference bits of each page



## ■ Second–chance algorithm

- Basically, FIFO algorithm
- If reference bit of the chosen page is 1, give a second chance
  - The reference bit is cleared.



# Page-Buffering Algorithms



## ■ Pool of free frames

- Some system maintain a list of free frames.
- When a page fault occurs, a victim frame is chosen from the pool.
- Frame number of each free page can be kept for next use.

## ■ Keep a list of free frames and remember which page was in each frame.

- The old page can be reused.
- OS typically applies ZFOD(zero-fill on demand) technique.

## ■ List of modified pages

- Whenever the page device is idle, a modified page is written to disk.

# Allocation of Frames



- How do we allocate the fixed amount of free memory among various processes?
  - How many frames does each process get?
- Minimum number of frames for each process
  - # of frames for each process decreases
    - page-fault rate is increases
    - performance degradation
  - Minimum # of frames should be large enough to hold all different pages that any single instruction can reference.

# Allocation Algorithms

## ■ Equal allocation

- Split  $m$  frames among  $n$  processes  
→  $m/n$  frames for each process

## ■ Proportional allocation 크기의 비례하여 구분 방법

- Allocate available memory to each process according to its size

$$a_i = s_i / S * m$$

- $a_i$ : # of frames allocated to process  $p_i$
  - $s_i$ : size of process  $p_i$
  - $S = \sum s_i$
- Variation: frame allocation based on ...
    - Priority of process
    - Combination of size and priority

# Global vs. Local Allocation

process의 경계 생각 X, 더 효율적인 선택 가능, 남의 공간도 나눠 쓸 수 있음

- **Global replacement**: a process can select a replacement frame from the set of all frames, including frames allocated to other processes

- A process cannot control its own page-fault rate

내꺼 공간이 바뀌고

- **Local replacement**: # of frames for a process does not change
- Less used pages of memory can't be used by other process

→ global replacement is more common method.



# Agenda

---



- Background
- Demand paging
- Copy-on-Write (COW)
- Page replacement
- Allocation of frames
- Thrashing
- Memory-mapped files
- Allocating kernel memory
- Other considerations
- Operation-system examples

# Thrashing




---

- If a process does not have enough frames to support pages in active use, it quickly faults again and again.
- A process is **thrashing** if it is spending more time paging than executing.

# Thrashing → process 별로 Thrashing이 안 일어나는 만큼 메모리를 주어야함.

- If thrashing sets in, CPU utilization drops sharply.
  - Degree of multiprogramming should be decreased



# Thrashing

memory access가 균일하게 일어나지 않는다.

시간에 따라 기억된 데이터의 접근 패턴이 많이 다르다.

access 위해 page마다 같이 쓰는 페이지가 많다.

불가피한 대수식 > frame 수  
→ page fault rate이 급격하게 증가


조금씩이 memory access 증가

- To prevent thrashing, a process must be provided with as many frames as it needs.

→ How to know how many frames it needs?

## ■ Locality model

- Locality: set of pages actively used together
- A program is generally composed of several localities




# Working-Set Model

- **Working set**: set of pages in the most recent  $\Delta$  page references
  - Parameter  $\Delta$ : **working-set window**

page reference table


... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



- $WSS_i$ : **working set size** of process  $p_i$
- **Process  $p_i$  needs  $WSS_i$  frames**
- If total demand is greater than # of available frames, thrashing will occur.


# Working Sets and Page Fault Rates

- Direct relationship between working set of a process and its page-fault rate
- Working set changes over time
- Peaks and valleys over time



# Page-Fault Frequency

- Alternative method to control trashing: control degree of multiprogramming by page-fault frequency (PFF)
  - If PFF of a process is too high, allocate more frame
  - If PFF of a process is too low, remove a frame from it



# Agenda

---



- Background
- Demand paging
- Copy-on-Write (COW)
- Page replacement
- Allocation of frames
- Thrashing
- **Memory-mapped files**
- Allocating kernel memory
- Other considerations
- Operation-system examples




# Memory-Mapped Files



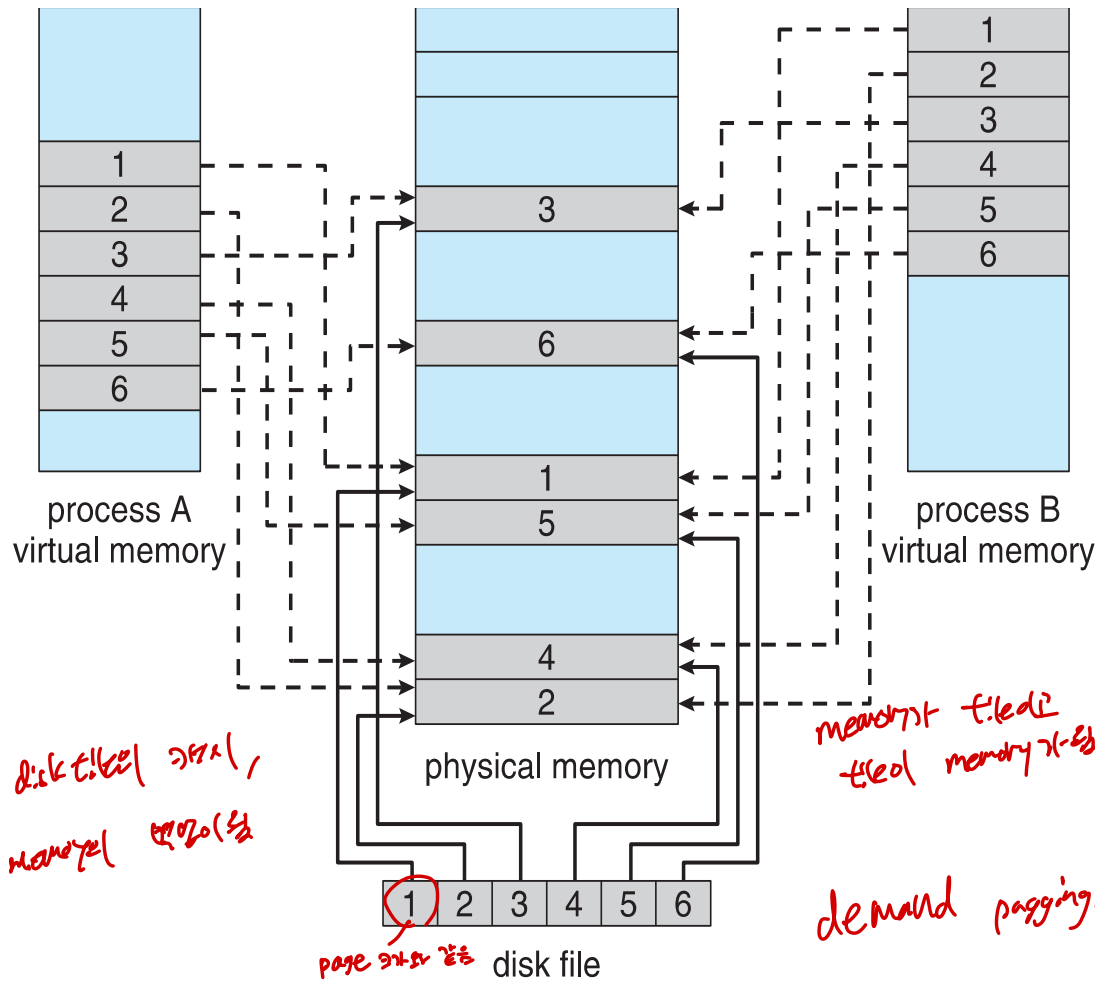
- **Memory-mapped file:** a part of virtual address space is logically associated with a file
  - Map a disk block to a page in memory
    1. Initial access proceeds through demand paging results in page fault.
    2. A page-sized portion of the file is read from the file system into a physical page.
    3. Subsequent access is through memory access routine.
    4. When the file is closed, memory-mapped data are written back to disk.
- **Advantages**
  - Reduces overhead of read() and write()
  - File sharing

# Memory-Mapped Files



# Memory-Mapped Files

demand paging의 공유하는 일대일 사용 가능



physical memory disk에 있는 것,  
disk에 physical memory 반영이 됨


memory가 tiled된  
tiled memory 구조

demand paging과 같음

Memory-mapped file  
Shared by two processes

# Shared Memory in Win32 API

- Sharing memory-mapped file is similar to shared memory
  - On POSIX and Windows, shared memory is accomplished by memory mapping files




# Shared Memory in Win32 API

## ■ Process 1

1. Create a file to be shared
  - `CreateFile(...)`
2. Create file mapping (named shared memory object)
  - `CreateFileMapping(...)`
3. Establish a view of mapped file in virtual address space
  - `MapViewOfFile(...)`


## ■ Process 2

1. Open file mapping
  - `OpenFileMapping(...)`
2. Establish a view of mapped file in virtual address space
  - `MapViewOfFile(...)`



# Memory Mapped File on Windows

- File mapping is the association of a file's contents with a portion of the virtual address space of a process.
  - The system creates a **file mapping** object to maintain this association.
  - A **file view** is the portion of virtual address space that a process uses to access the file's contents.
  - It also allows the process to work efficiently with a large data file.
  - Multiple processes can also use memory-mapped files to share data.



# Reading Assignment


---



- Visit and study the following web pages to learn memory mapped file on Windows
  - File mapping
    - [https://msdn.microsoft.com/en-us/library/windows/desktop/aa366556\(v=vs.85\).aspx](https://msdn.microsoft.com/en-us/library/windows/desktop/aa366556(v=vs.85).aspx)
  - File mapping functions
    - [https://msdn.microsoft.com/en-us/library/windows/desktop/aa366781\(v=vs.85\).aspx#file\\_mapping\\_functions](https://msdn.microsoft.com/en-us/library/windows/desktop/aa366781(v=vs.85).aspx#file_mapping_functions)

# Memory-Mapped I/O

- I/O devices are accessed through ...
  - Device registers in I/O controller to hold command and data
  - Usually, special instructions transfer data between device registers and memory





# Memory-Mapped I/O



- **Memory-mapped I/O:** ranges of memory addresses are set aside and mapped to the device registers
    - In IBM PC, each location on screen is mapped to a memory location
    - Serial/parallel ports – data transfer through reading/writing device registers (ports)
  
  - **Programmed I/O (PIO)**
    - Ex) Sending a long string of bytes through memory-mapped I/O
      - CPU sets a byte to a register and set a bit in the control register to signal that the data is available.
      - Device receives the data and clears the bit in the control register to signal CPU that it is ready for the next byte.
- cf. Interrupt driven I/O

# Agenda

---



- Background
- Demand paging
- Copy-on-Write (COW)
- Page replacement
- Allocation of frames
- Thrashing
- Memory-mapped files
- Allocating kernel memory paging보다 더한게 쉬운 것
- Other considerations
- Operation-system examples

# Allocating Kernel Memory




- Allocation of kernel memory requires special handling
  - Kernel requests memory for data structures of varying sizes
    - Many OS's do not subject kernel code/data to the paging system
  - Certain H/W devices interact directly with physical memory
    - Memory should reside in physically contiguous pages.
- Strategies for kernel memory allocation
  - Buddy system
  - Slab allocation

# Buddy System

- **Buddy system**: allocates memory from a fixed-size segment consisting of physically contiguous pages

- Power-of-2 allocator

Ex) initially 256 KB is available, 21 KB was requested



- Advantage: easy to combine adjacent buddies
- Disadvantage: internal fragmentation  $12 < 32$

frac 할때 2의 제곱

# Slab Allocation


structure 포인터로 -> 크기 일정  
ex) PCB, TCB, shared memory block

- Motivation: mismatch between allocation size and requested size
  - Page-size granularity vs. byte-size granularity
  - Applied since Solaris 2.4 and Linux 2.2

- Cache for each unique kernel data structure

동일한 크기의 structure가 반복 사용됨

- A **slab** is made up of one or more physically contiguous pages
- A **cache** consists of one or more slabs



# Slab Allocation




- Single cache is for each unique kernel data structure
  - Each cache filled with **objects** – instantiations of the data structure.  
Ex) cache for process descriptor, cache for file objects, cache for semaphore, ...
- When cache created, filled with objects marked as **free**.
- When structures stored, objects marked as **used**.
- If slab is full of used objects, next object allocated from empty slab.
  - If no empty slabs, new slab is allocated.

# Slab Allocation

## ■ Benefits

- No memory waste due to fragmentation
- Memory requests can be satisfied quickly
- ➔ Suitable for data structures that are allocated and deallocated frequently.



# Other Considerations

---



- Prepaging
- Page size
- TLB reach
- Inverted page tables
- Program structure
- I/O interlock



# Prepaging

page faults를 줄이려는 것  
↓  
다시 예측하기

- A problem of pure demand paging: a large number of page faults
- **Prepaging**: bring all pages that will be needed at one time to reduce page faults.
  - Ex) working-set model (같이 다비는 page fault set)
  - Important issue: cost of prepaging vs. cost of servicing corresponding page faults

# Page Size 증가하는 추세

intel page size { 4k  
4M

## ■ Issues about page size

	smaller page	larger page
Size of page table	large	small
Memory utilization	better	worse
I/O latency	large	small
Locality	good	bad
Page fault	many	few

memory overhead

page fault rate

한번 더 cost가 큼

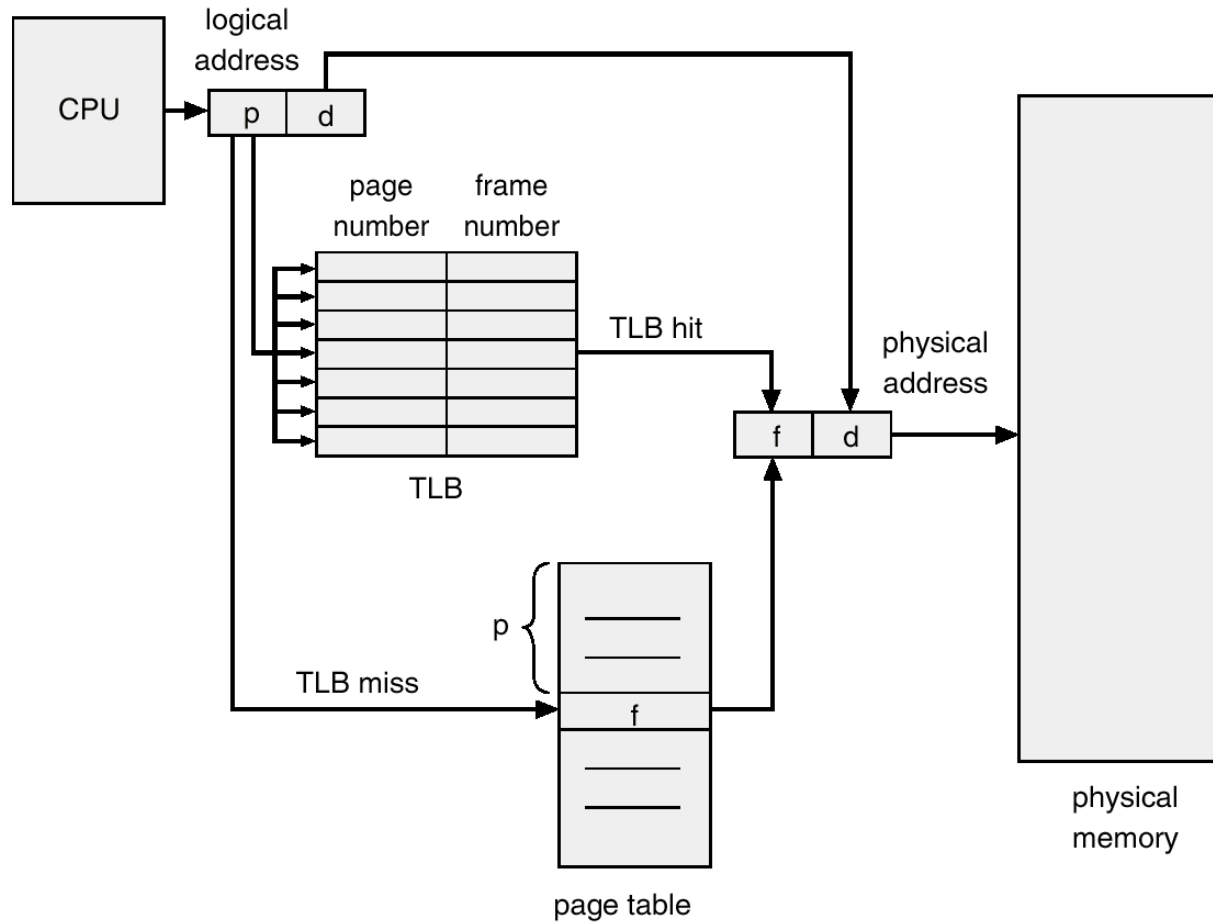
## ■ Historical trend: page size is getting larger

# TLB Reach



- To improve **TLB hit ratio**, size of TLB should be increased.  
→ but associate memory is expensive, power hungry
- **TLB reach**: amount of memory accessible from TLB
  - $\text{TLB reach} = \langle \# \text{ of entries in TLB} \rangle * \langle \text{page size} \rangle$
- TLB reach can increase by increasing page size
  - However, with large page, fragmentation also increases.  
→ S/W managed TLB (OS support several different page sizes)  
Ex) UltraSparc, MIPS, Alpha  
Cf) PowerPC, Pentium: H/W managed TLB

# TLB Reach



# Program Structure

- User don't have to know about nature of memory. But, if user knows underlying demand paging, performance can be improved

Ex) If page size is 128 words, B is better than A

```
for(j = 0; j < 128; j++)  
    for(i = 0; i < 128; i++)  
        data[i][j] = 0;
```


} A

```
for(i = 0; i < 128; i++)  
    for(j = 0; j < 128; j++)  
        data[i][j] = 0;
```

} B

# I/O Interlock

- Memory space related to I/O transfer should not be replaced
- Remedies
  - I/O transfer is performed only through system memory
  - Locking pages in memory



# Agenda

---



- Background
- Demand paging
- Copy-on-Write (COW)
- Page replacement
- Allocation of frames
- Thrashing
- Memory-mapped files
- Allocating kernel memory
- Other considerations
- Operation-system examples



## ■ Demand paging


- Allocates pages from a list of free frames.
- A global page-replacement policy similar to the LRU-approximation clock algorithm

## ■ Accessed bit

- Each page has an accessed bit that is set whenever the page is referenced.
- Periodically, the accessed bits for pages in the active list are reset

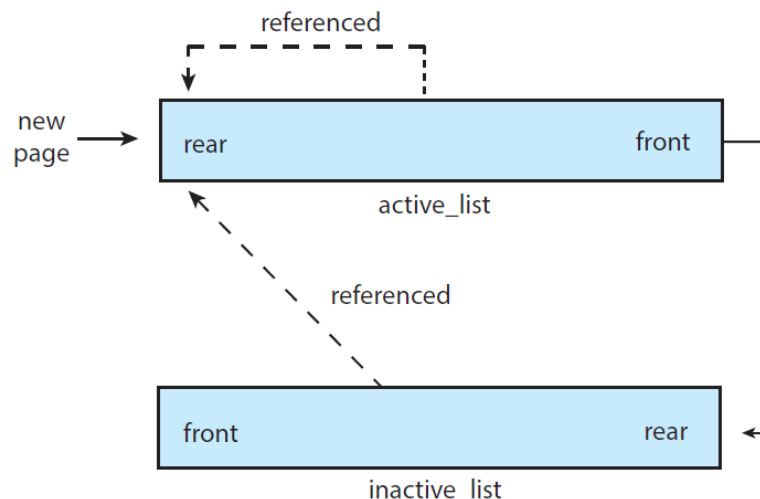


- Linux maintains two types of page lists
  - **Active list** contains the pages that are considered in use
    - Over time, the least recently used page will be at the front of the active list.
  - **Inactive list** contains pages that have not recently been referenced and are eligible to be reclaimed.
    - If a page in the inactive list is referenced, it moves back to the rear of the active list



## ■ Kernel swap daemon process kswapd

- A page-out daemon process
- Periodically awakens and checks the amount of free memory in the system
- If free memory falls below a certain threshold, kswapd begins scanning pages in the inactive list and reclaiming them for the free list



# Windows



- Demand paging with **clustering**
  - At page fault, Windows brings not only the fault page but also several following pages
- Working-set minimum/maximum
  - **Working-set minimum** (default=50): the virtual memory manager attempts to keep **at least this much memory resident** in the process whenever the process is active
  - **Working-set maximum** (default: 345): the virtual memory manager attempts to keep **no more than this much memory resident** in the process whenever the process is active and available memory is low.